



Chapter 17: Indexing Structures

CS-6360 Database Design

Chris Irwin Davis, Ph.D.

Email: cid021000@utdallas.edu

Phone: (972) 883-3574

Office: ECSS 4.705

Outline

- Types of Single-level Ordered Indexes
 - Primary Indexes
 - Clustering Indexes
 - Secondary Indexes
- Multilevel Indexes
- Dynamic Multilevel Indexes Using B-Trees and B⁺-Trees

- A single-level index is an auxiliary file that makes it more efficient to search for a record in the data file.
- In a database, the index is usually stored in the same file as the data.
- The index is usually specified on one field of the file (although it could be specified on several fields)
- One form of an index is a file of entries **<field value, pointer to record>**, which is ordered by field value

- The index is called an *access path* on the field.
- The index file usually occupies considerably less disk blocks than the data file because its entries are much smaller
- A binary search on the index yields a pointer to the file record
- Indexes can also be characterized as dense or sparse
 - A **dense index** has an index entry for every search key value (and hence every record) in the data file.
 - A **sparse (or nondense) index**, on the other hand, has index entries for only some of the search values

Single-level Ordered Indexes

Primary Index

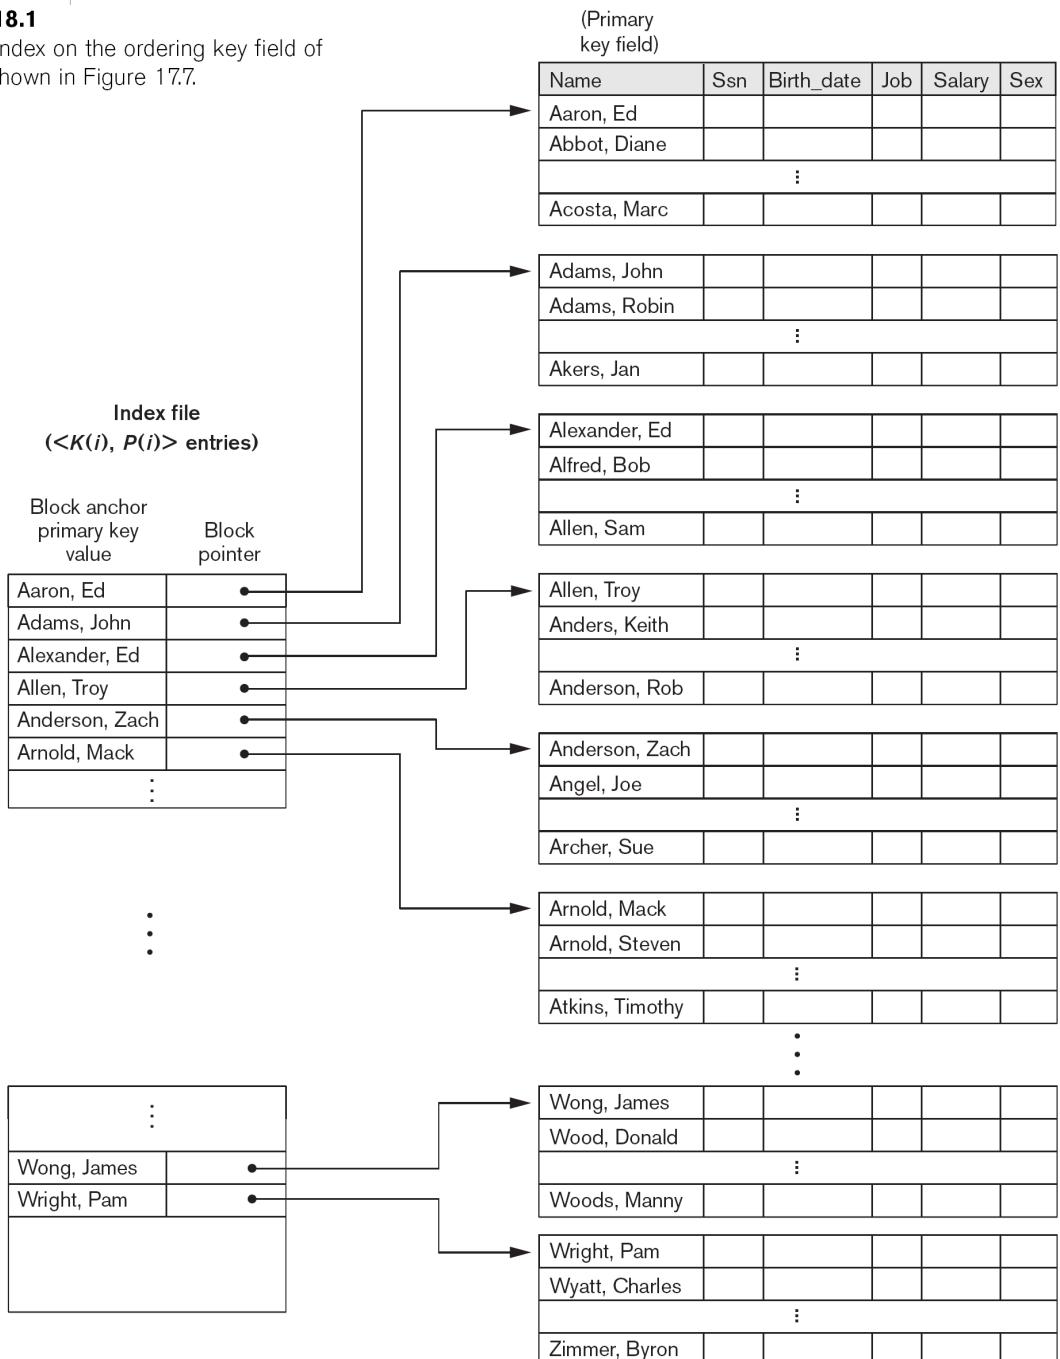
Types of Single-Level Indexes

- Primary Index
 - Defined on an ordered data file
 - The data file is ordered on a **key field**
 - Includes one index entry *for each block* in the data file; the index entry has the key field value for the *first record* in the block, which is called the *block anchor*
 - A similar scheme can instead use the *last record* in a block.
 - A primary index is a nondense (sparse) index, since it includes an entry for each disk block of the data file and the keys of its anchor record rather than for every search value.

Primary Index on the Ordering Key Field

Figure 18.1

Primary index on the ordering key field of the file shown in Figure 17.7.



Single-level Ordered Indexes

Clustering Index

- Clustering Index
 - Defined on an ordered data file
 - The data file is ordered on a *non-key field* unlike primary index, which requires that the ordering field of the data file have a distinct value for each record.
 - Includes one index entry *for each distinct value* of the field; the index entry points to the first data block that contains records with that field value.
 - It is another example of *nondense* index where Insertion and Deletion is relatively straightforward with a clustering index.

A Clustering Index Example

- Clusters may begin in the middle of a block
- Clusters span blocks
- No empty block slots
- Bad for dynamic size

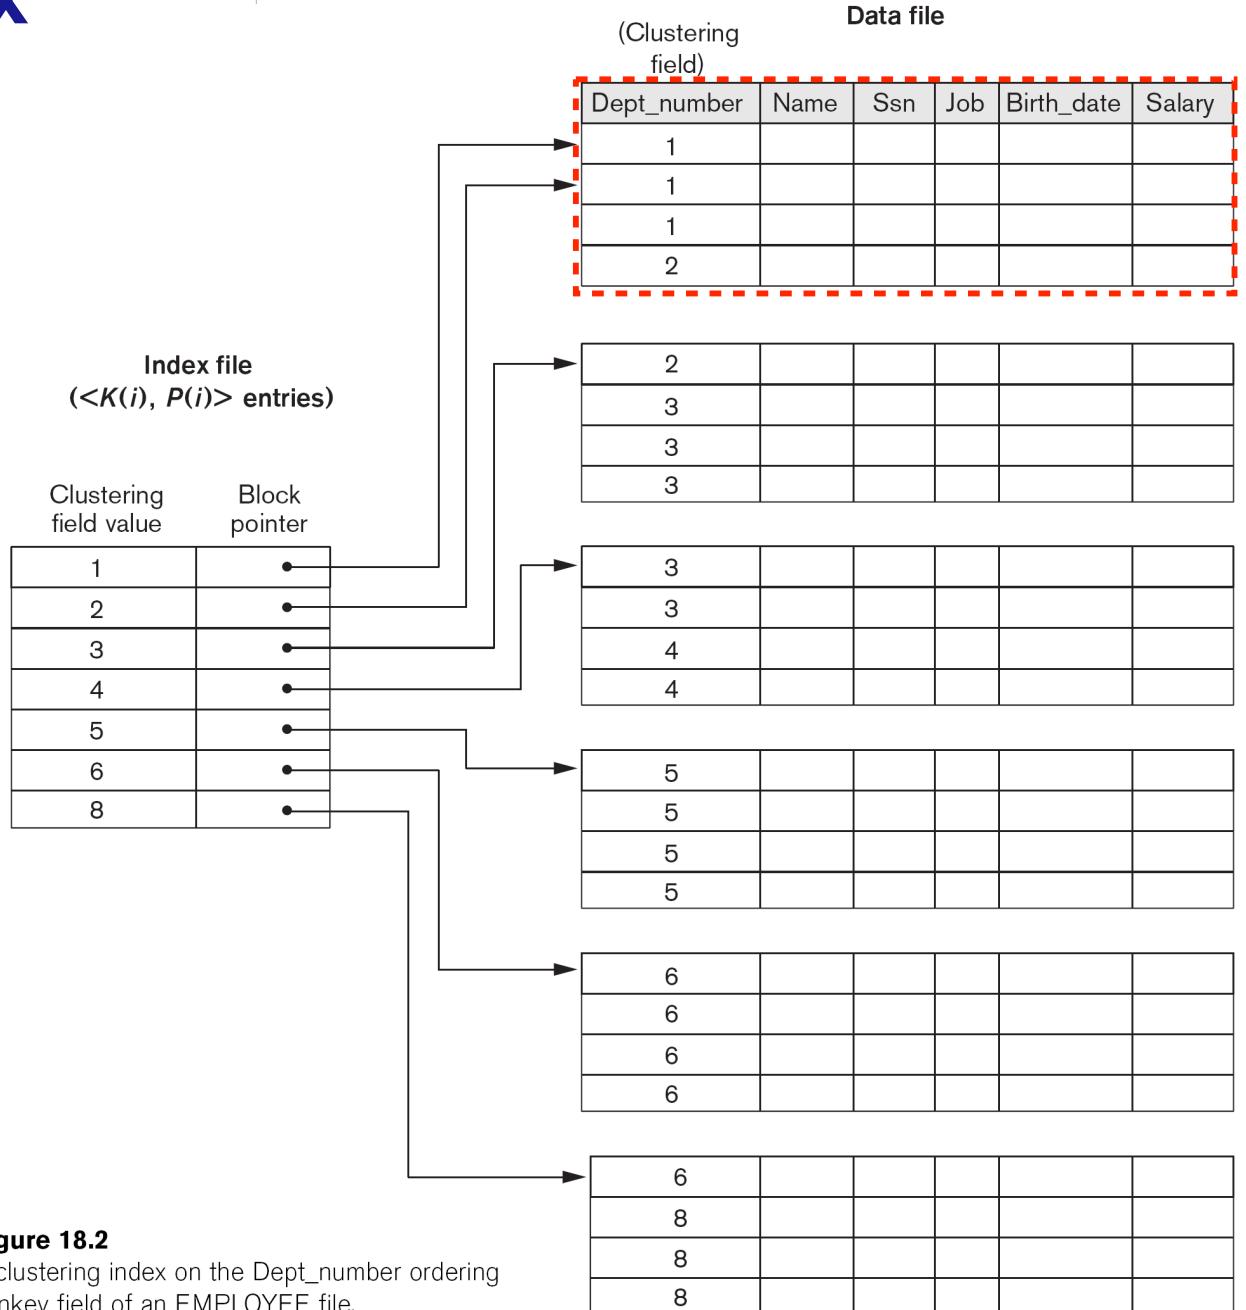


Figure 18.2

A clustering index on the Dept_number ordering nonkey field of an EMPLOYEE file.

A Clustering Index Example

- Clusters may begin in the middle of a block
- Clusters span blocks
- No empty block slots
- Bad for dynamic size

Pointing to the whole block, not individual records

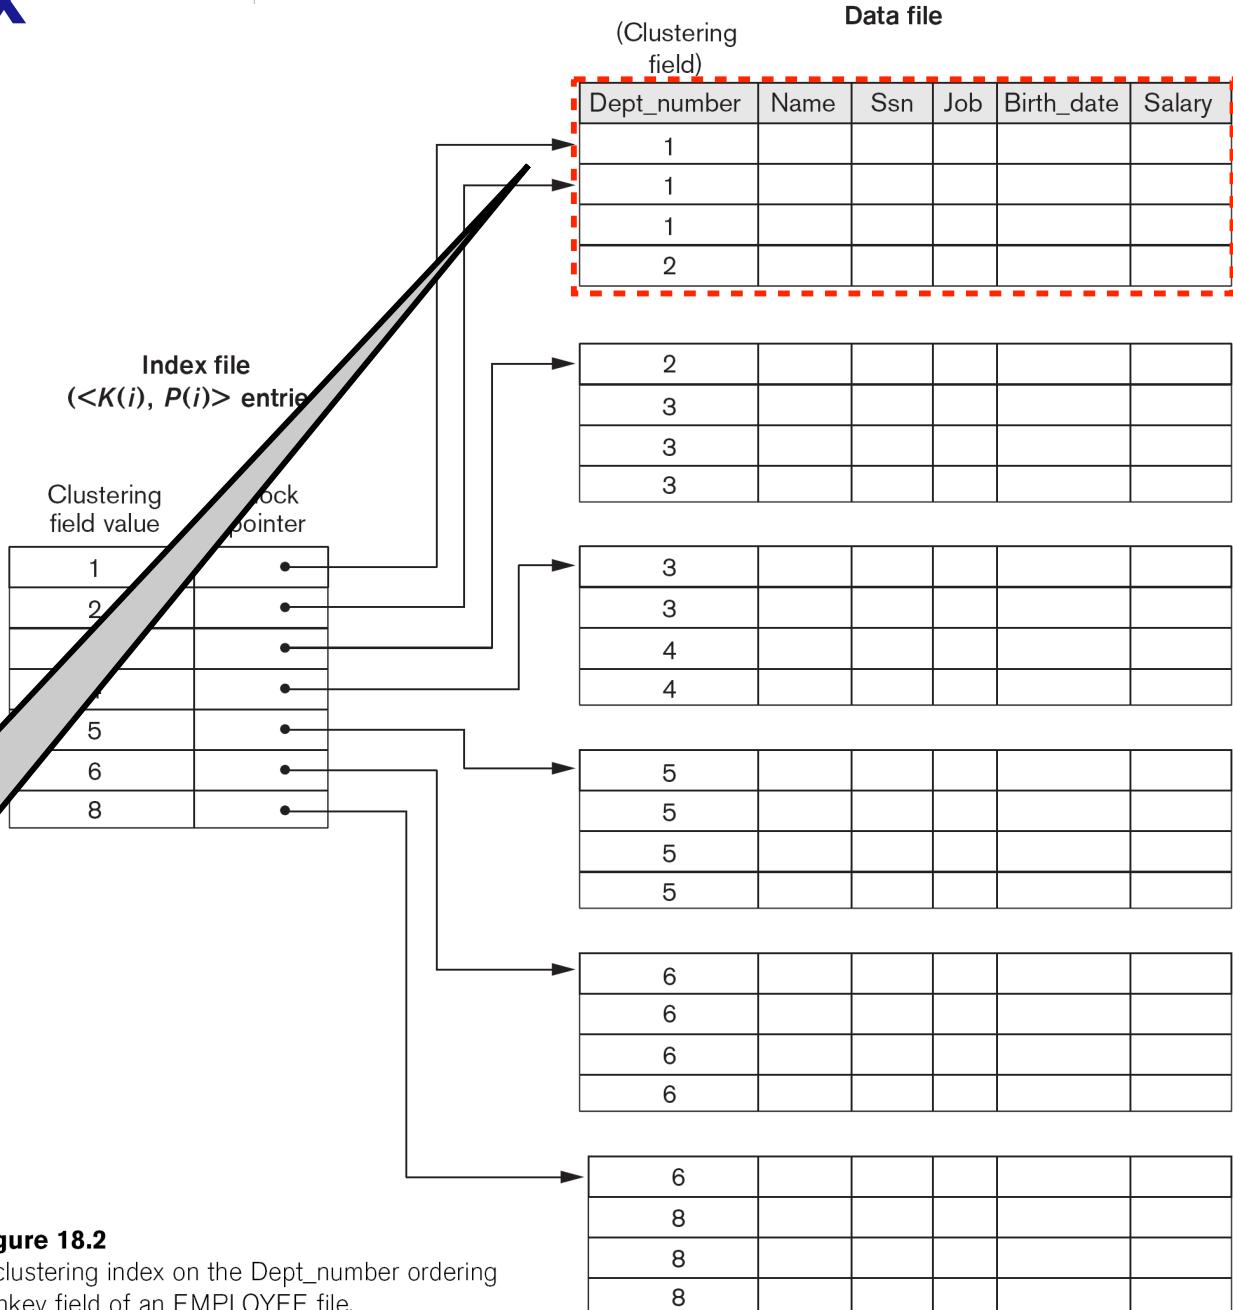


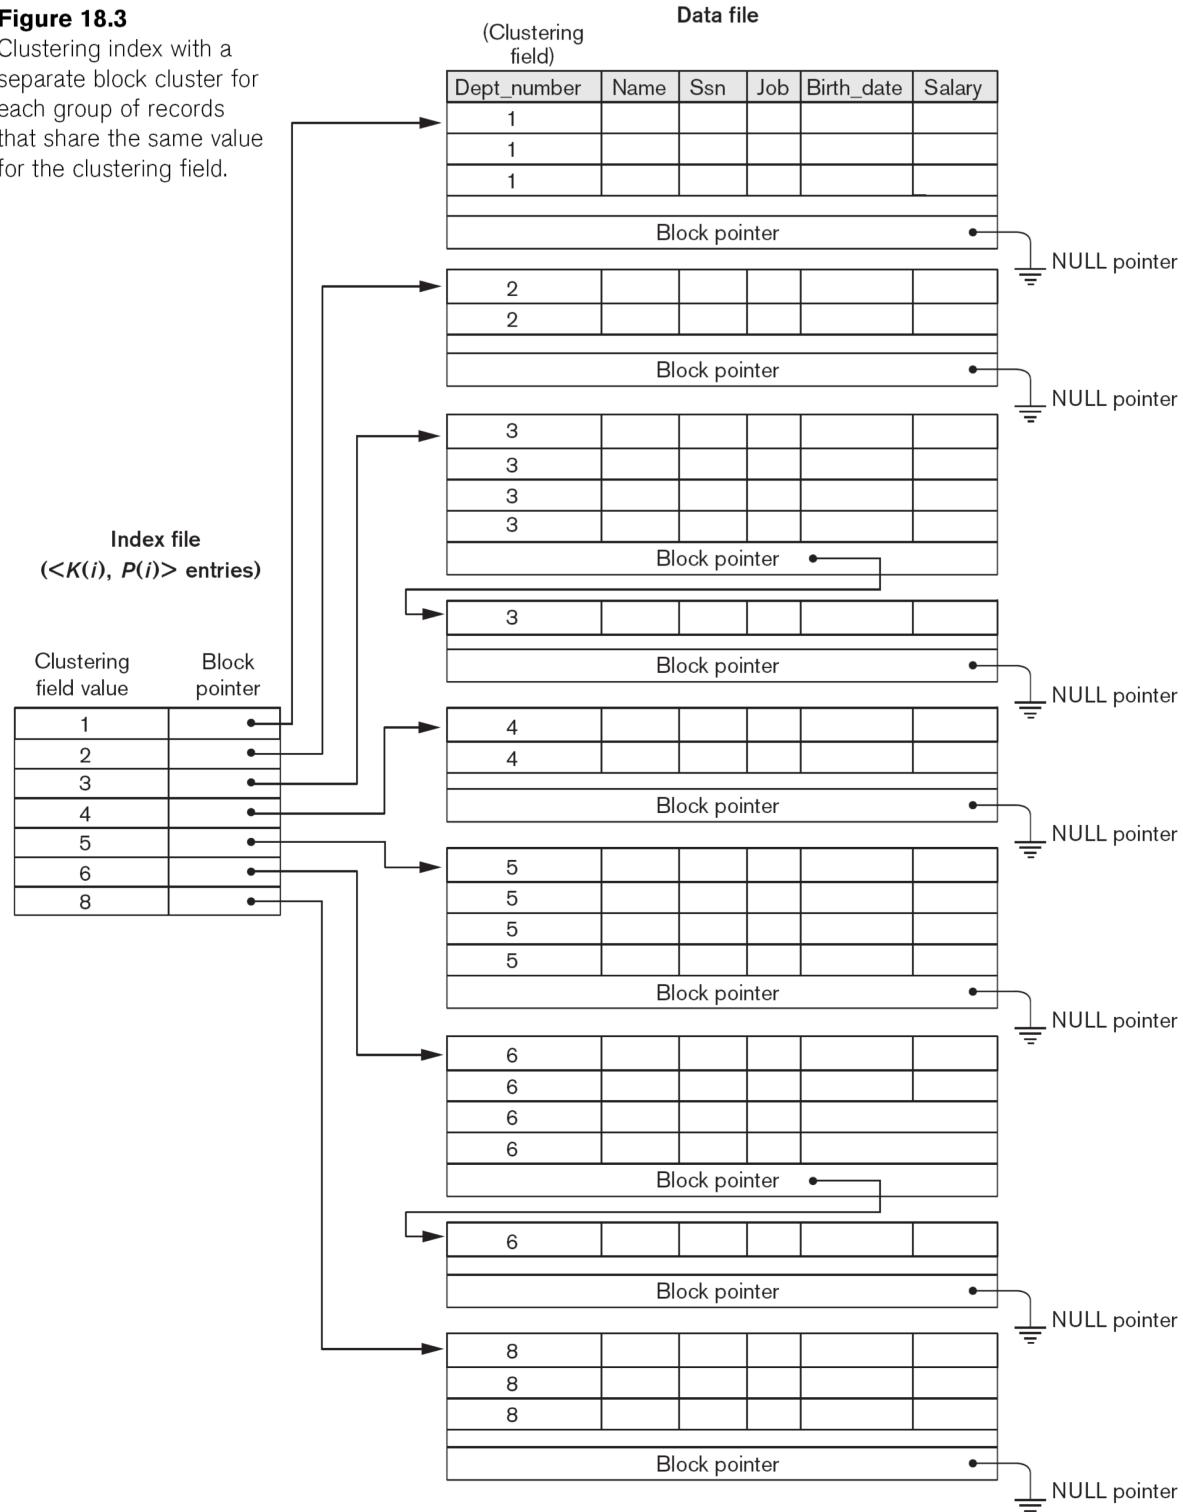
Figure 18.2

A clustering index on the Dept_number ordering nonkey field of an EMPLOYEE file.

Another Clustering Index Example

- No two clusters share the same block space
- Overflow block per cluster
- Multiple partially full blocks
- Accommodates dynamic size
- Lost performance due to redirection

Figure 18.3
Clustering index with a separate block cluster for each group of records that share the same value for the clustering field.



Single-level Ordered Indexes

Secondary Index

Types of Single-Level Indexes

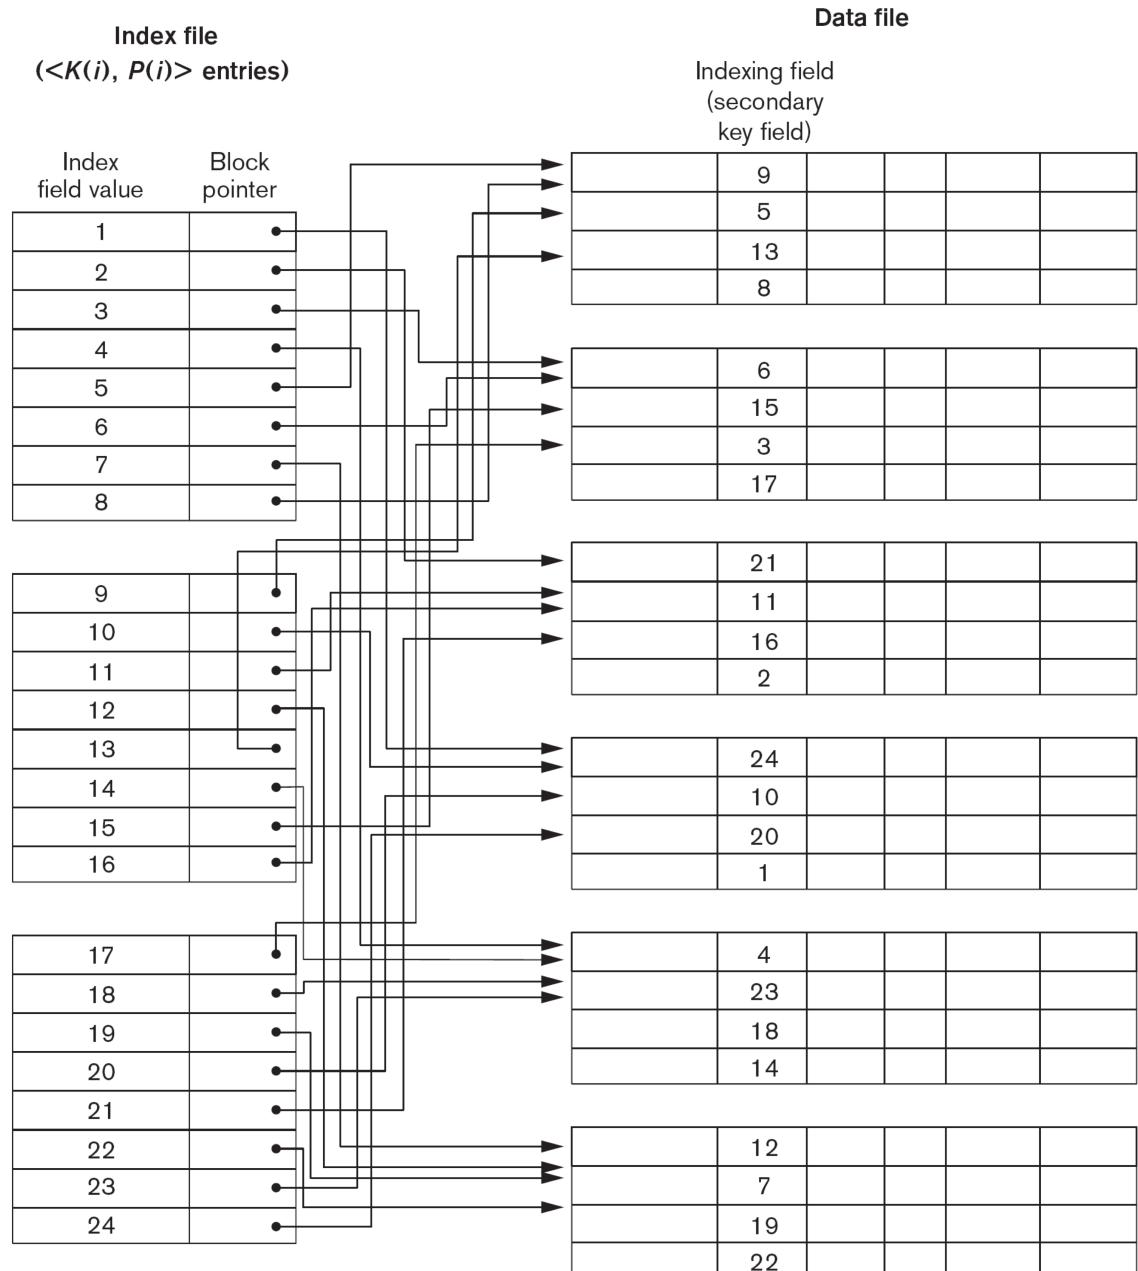
- Secondary Index
 - A secondary index provides a secondary means of accessing a file for which some primary access already exists.
 - The secondary index may be on a field which is a **candidate key** and has a unique value in every record, **OR** a **non-key** with duplicate values.
 - The index is an ordered file with two fields.
 - The first field is of the same data type as some **non-ordering field** of the data file that is an indexing field.
 - The second field is either a **block** pointer or a record pointer.
 - There can be *many* secondary indexes (and hence, indexing fields) for the same file.
 - Includes one entry for *each record* in the data file; hence, it is a *dense index*

Example of a Dense Secondary Index

- The entries are ordered by value of $K(i)$, so we can perform a binary search.
- Because the records of the data file are not physically ordered by values of the secondary key field, we cannot use block anchors (i.e. dense)

Figure 18.4

A dense secondary index (with block pointers) on a nonordering key field of a file.



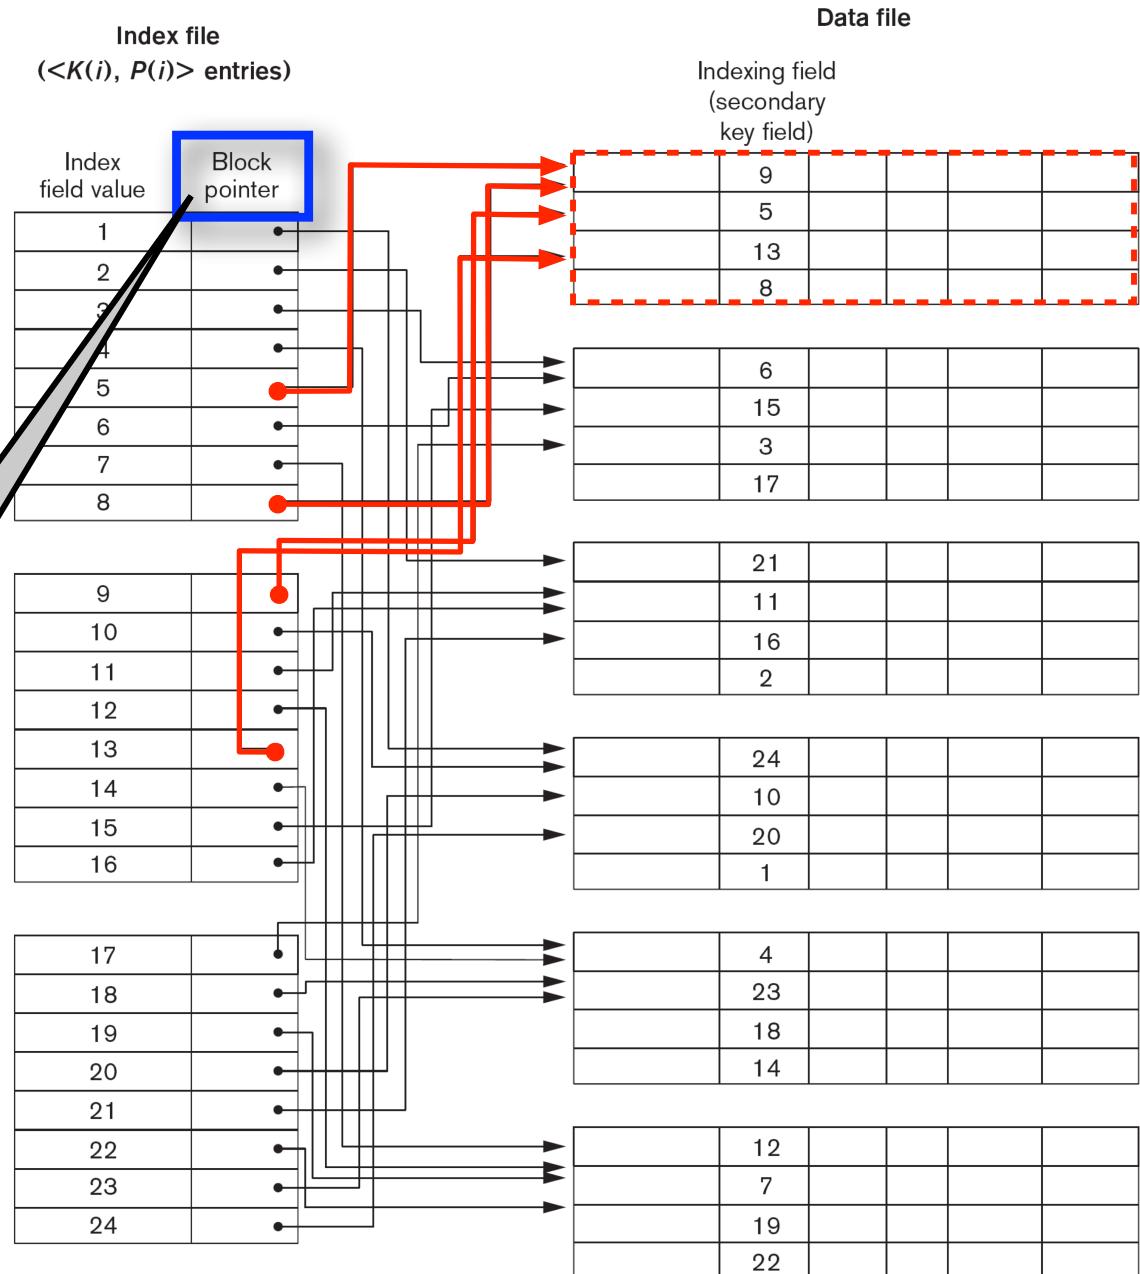
Example of a Dense Secondary Index

- The entries are ordered by value of $K(i)$, so we can perform a binary search.
- Because the records of the data file are not physically ordered by values of the secondary key field, we cannot use block anchors (i.e. dense)

Pointing to the whole block, not individual records

Figure 18.4

A dense secondary index (with block pointers) on a nonordering key field of a file.

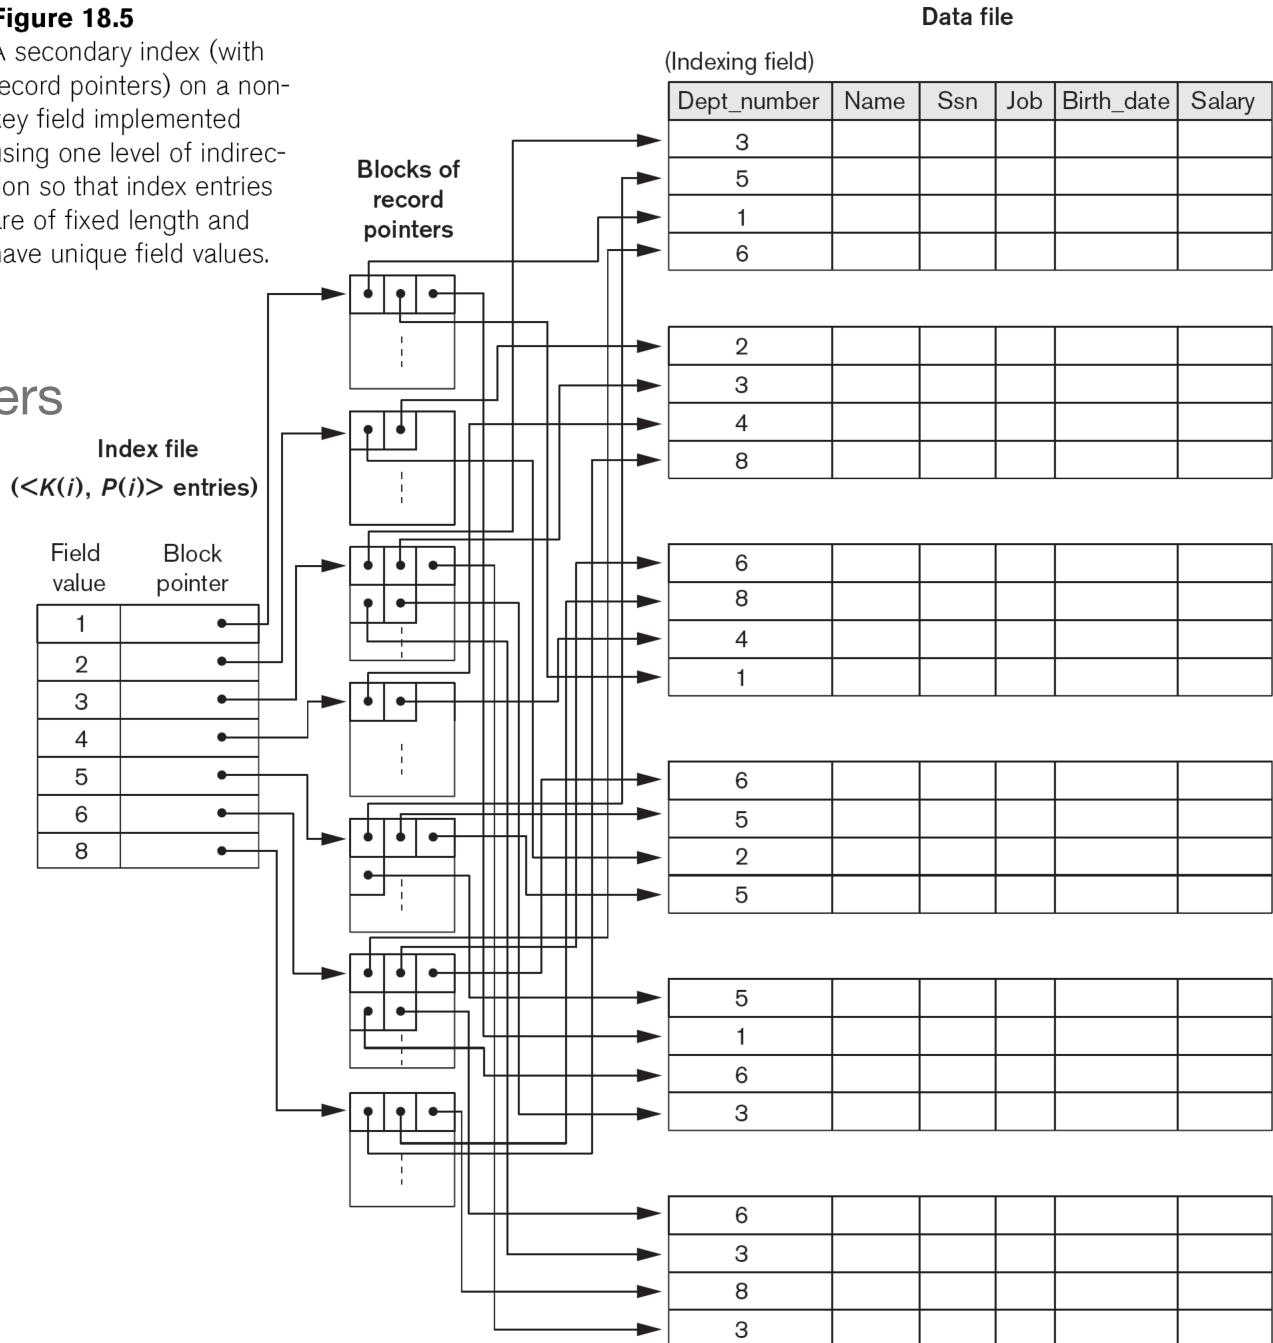


Example of a Secondary Index

Figure 18.5

A secondary index (with record pointers) on a non-key field implemented using one level of indirection so that index entries are of fixed length and have unique field values.

- Each block of record pointers may be a different size



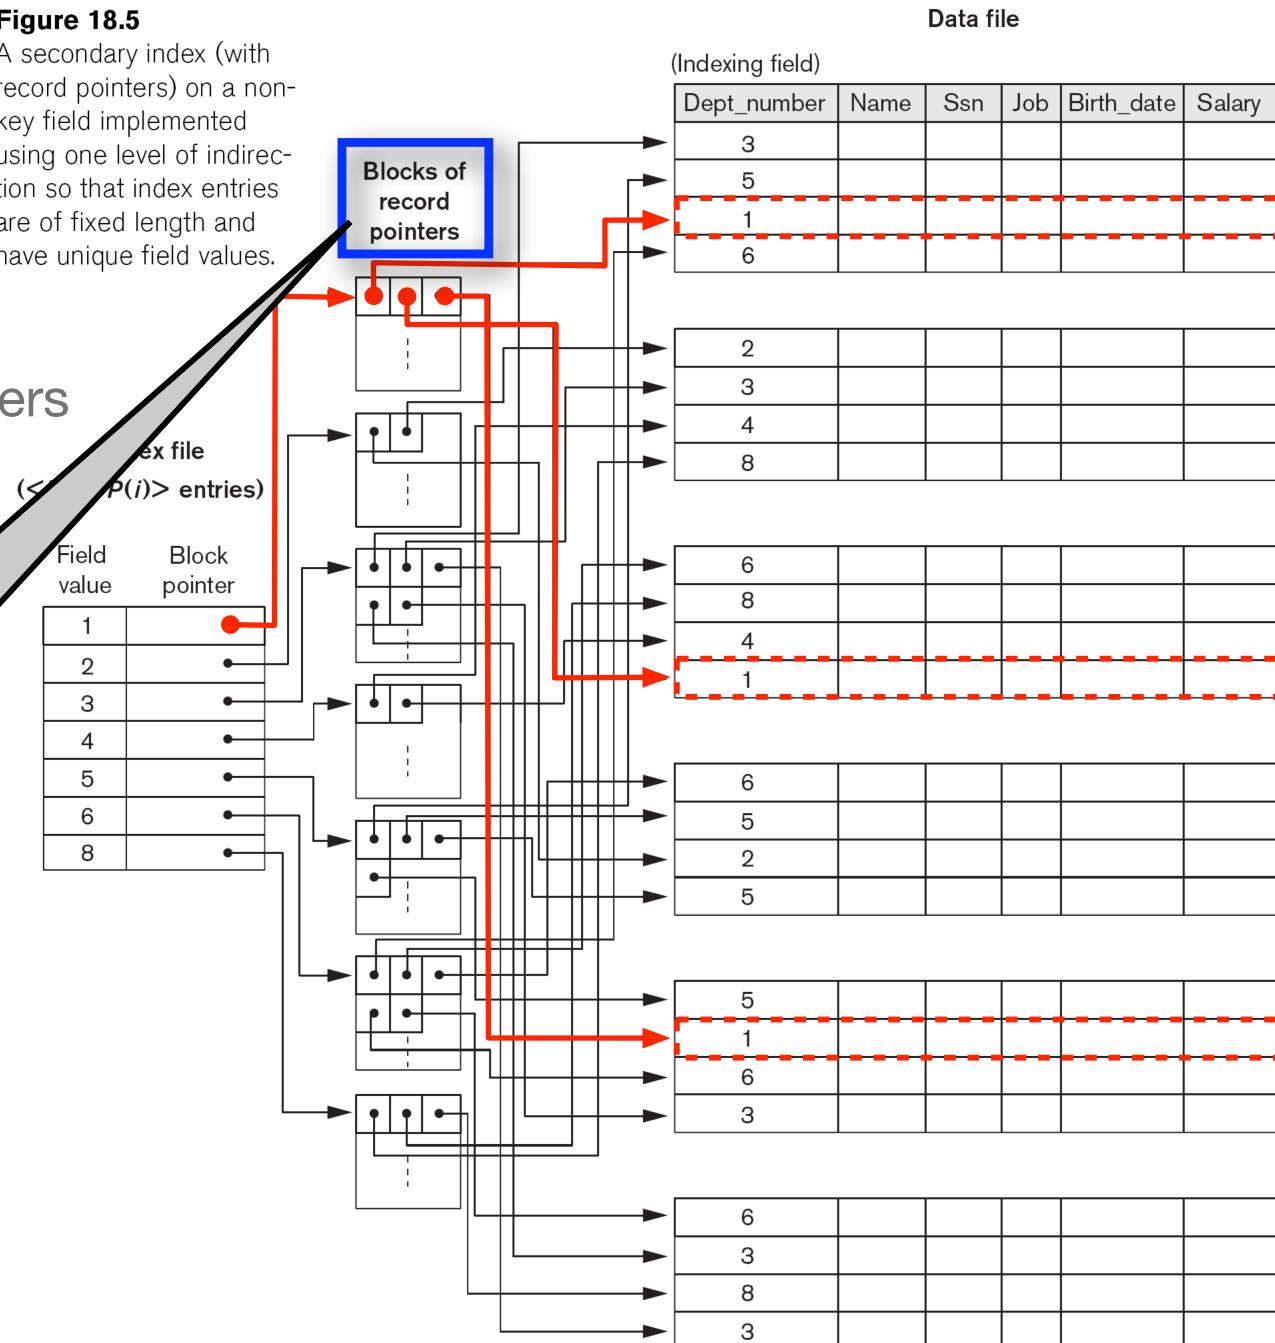
Example of a Secondary Index

Figure 18.5

A secondary index (with record pointers) on a non-key field implemented using one level of indirection so that index entries are of fixed length and have unique field values.

- Each block of record pointers may be a different size

Pointing to individual records



Properties of Index Types

Table 18.2 Properties of Index Types

Type of Index	Number of (First-level) Index Entries	Dense or Nondense (Sparse)	Block Anchoring on the Data File
Primary	Number of blocks in data file	Nondense	Yes
Clustering	Number of distinct index field values	Nondense	Yes/no ^a
Secondary (key)	Number of records in data file	Dense	No
Secondary (nonkey)	Number of records ^b or number of distinct index field values ^c	Dense or Nondense	No

^aYes if every distinct value of the ordering field starts a new block; no otherwise.

^bFor option 1.

^cFor options 2 and 3.

Multi-Level Indexes

Multi-Level Indexes

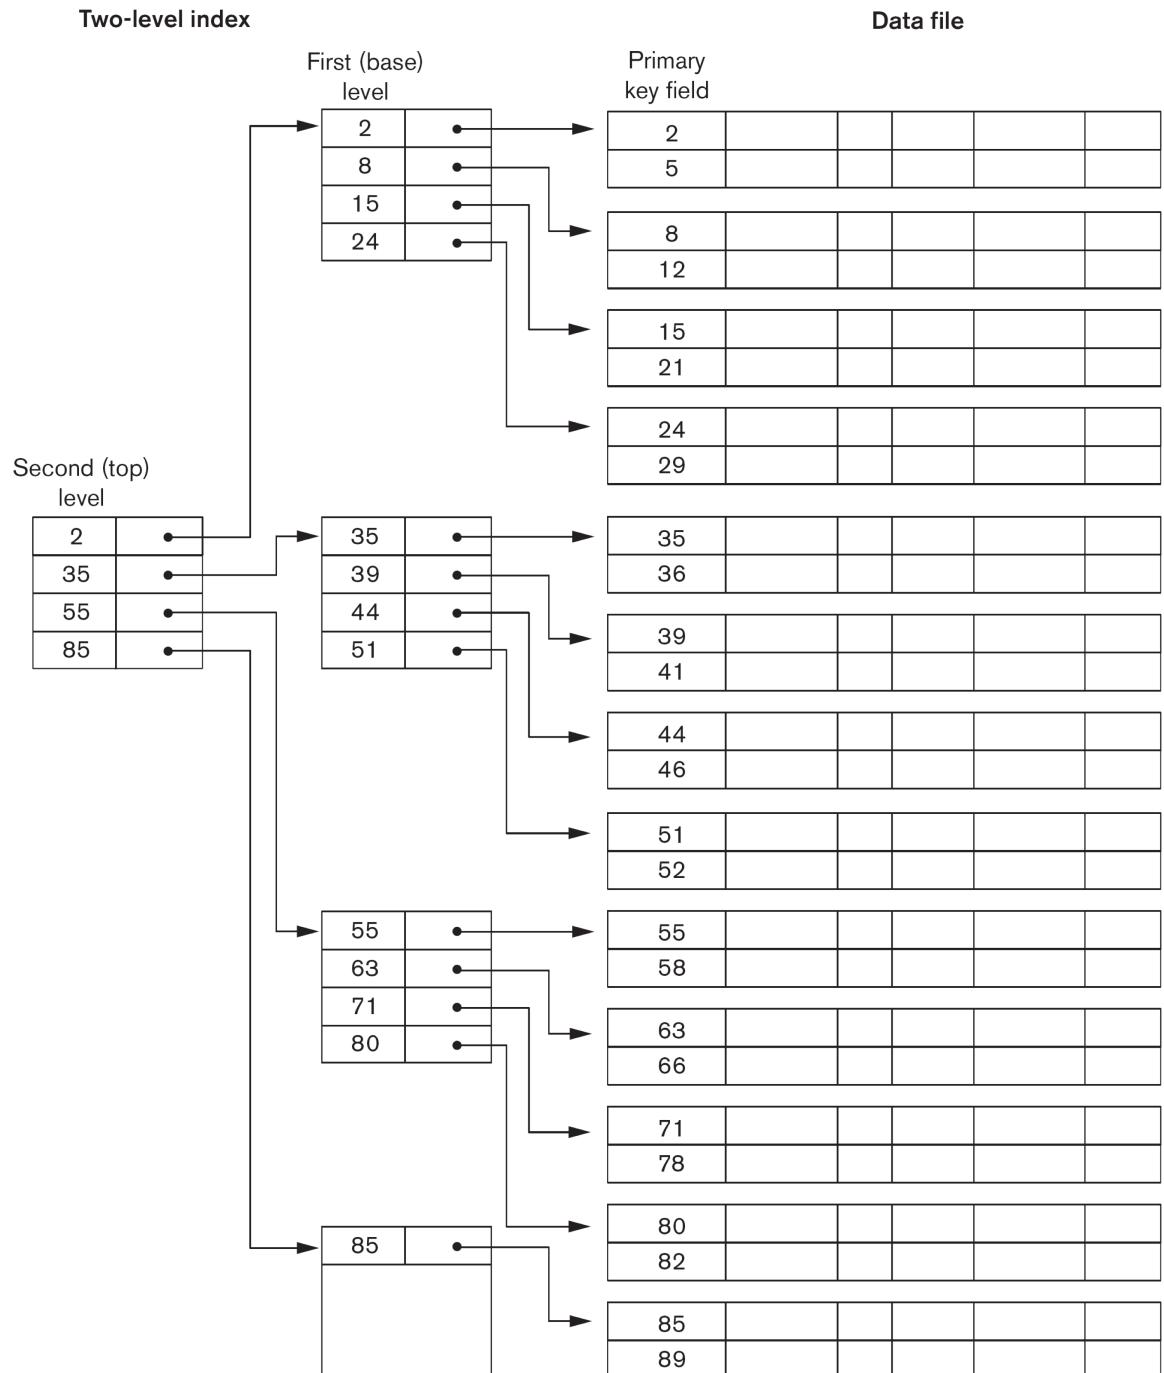
- Because a single-level index is an ordered file, we can create a primary index *to the index itself*;
 - In this case, the original index file is called the *first-level index* and the index to the index is called the *second-level index*.
- We can repeat the process, creating a third, fourth, ..., top level until all entries of the *top level* fit in one disk block
- A multi-level index can be created for any type of first-level index (primary, secondary, clustering) as long as the first-level index consists of *more than one* disk block

A Two-Level Primary Index

- Insertion is handled by some form of overflow file that is merged periodically with the data file.
- The index is recreated during file reorganization.

Figure 18.6

A two-level primary index resembling ISAM (Indexed Sequential Access Method) organization.

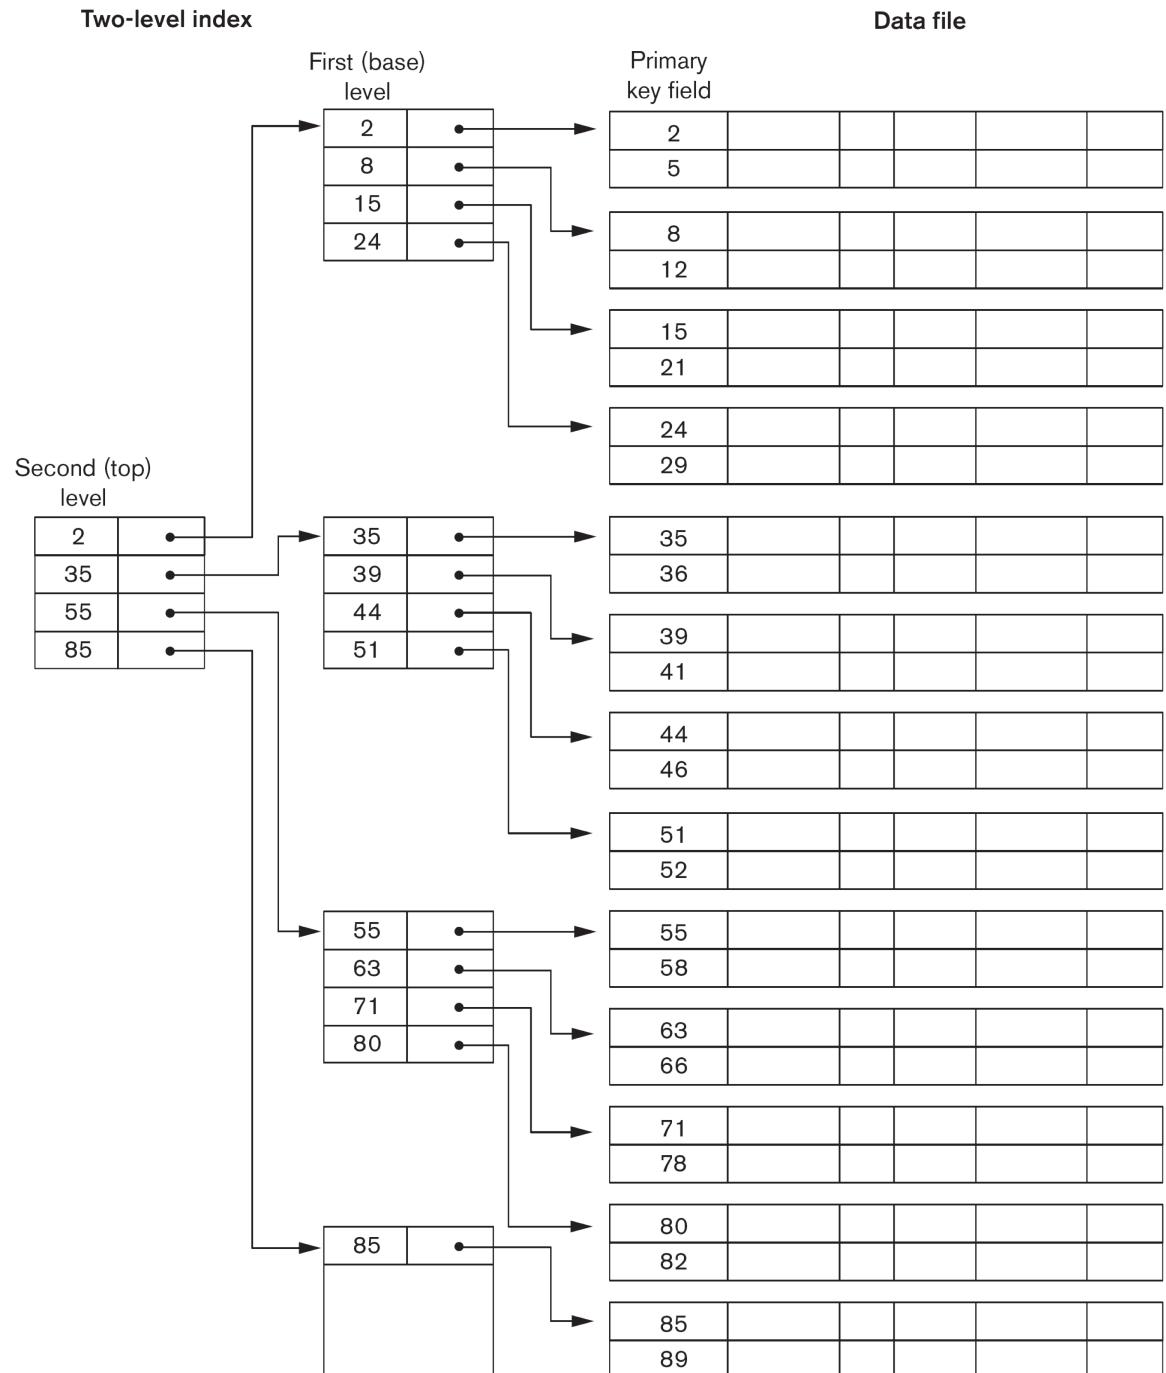


A Two-Level Primary Index

- Insertion and Deletion?
- Scalability?
- Flexibility?

Figure 18.6

A two-level primary index resembling ISAM (Indexed Sequential Access Method) organization.



Multi-Level Indexes

- Such a multi-level index is a form of *search tree*
 - However, insertion and deletion of new index entries is a severe problem because every level of the index is an *ordered file*.

- Such a multi-level index is a form of *search tree*
 - However, insertion and deletion of new index entries is a severe problem because every level of the index is an *ordered file*.

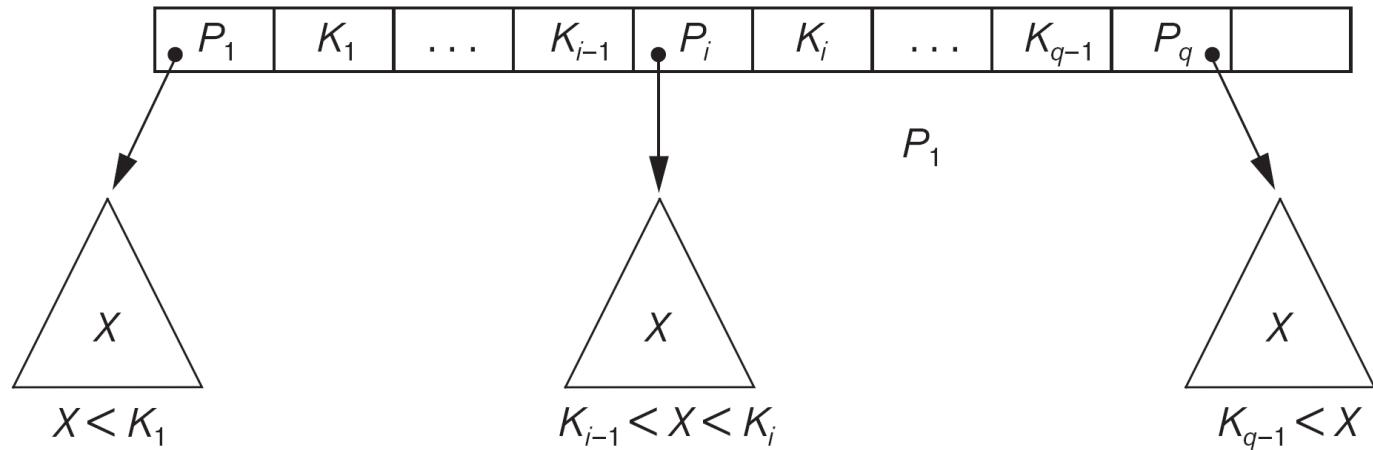
We can address this problem using...

Dynamic Multilevel Indexes Using B-Trees and B⁺-Trees

A Node in a Search Tree with Pointers to Subtrees Below It

Figure 18.8

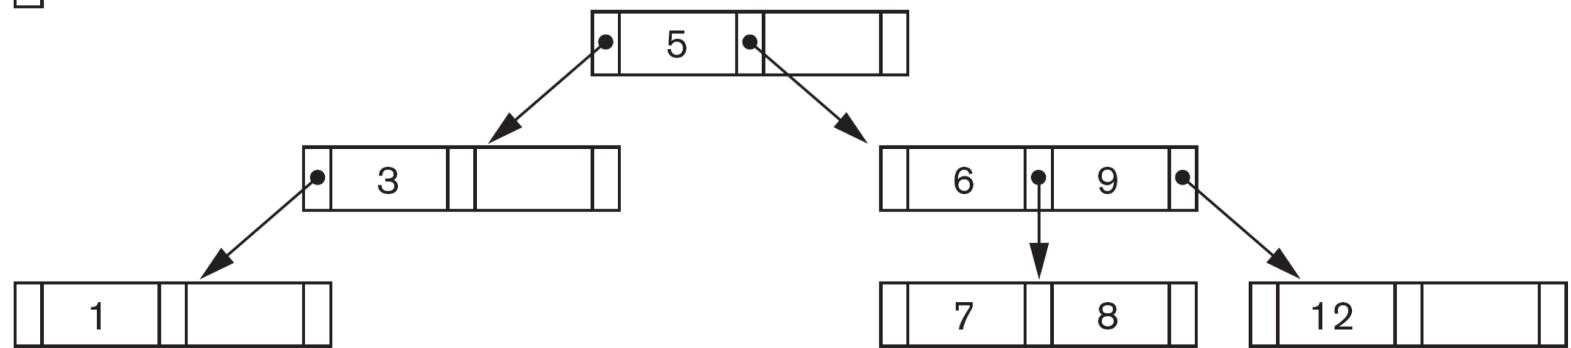
A node in a search tree with pointers to subtrees below it.



Node Detail

Figure 18.9
A search tree of
order $p = 3$.

- Tree node pointer
- Null tree pointer



Dynamic Multilevel Indexes Using B-Trees and B⁺-Trees



- Most multi-level indexes use B-tree or B⁺-tree data structures because of the insertion and deletion problem
 - This leaves space in each tree node (disk block) to allow for new index entries
- These data structures are variations of search trees that allow efficient insertion and deletion of new search values.
- In B-Tree and B⁺-Tree data structures, each node corresponds to a disk block
- Each node is kept between half-full and completely full

Dynamic Multilevel Indexes Using B-Trees and B⁺-Trees



- An insertion into a node that is not full is quite efficient
 - If a node is full, the insertion causes a split into two nodes
- Splitting may propagate to other tree levels
- A deletion is quite efficient if a node does not become less than half full
- If a deletion causes a node to become less than half full, it must be merged with neighboring nodes

Difference between B-tree and B⁺-tree

- In a B-tree, pointers to data records exist at **all levels** of the tree
- In a B⁺-tree, all pointers to data records exist at the **leaf-level** nodes
- A B⁺-tree can have fewer levels (or higher capacity of search values) than the corresponding B-tree

- In a B-Tree, half the time any search requires fewer comparisons
- In a B⁺ tree, key-sequential reads (i.e. **iteration**) can go across the bottom level

B-Tree Example

B-Tree Structures

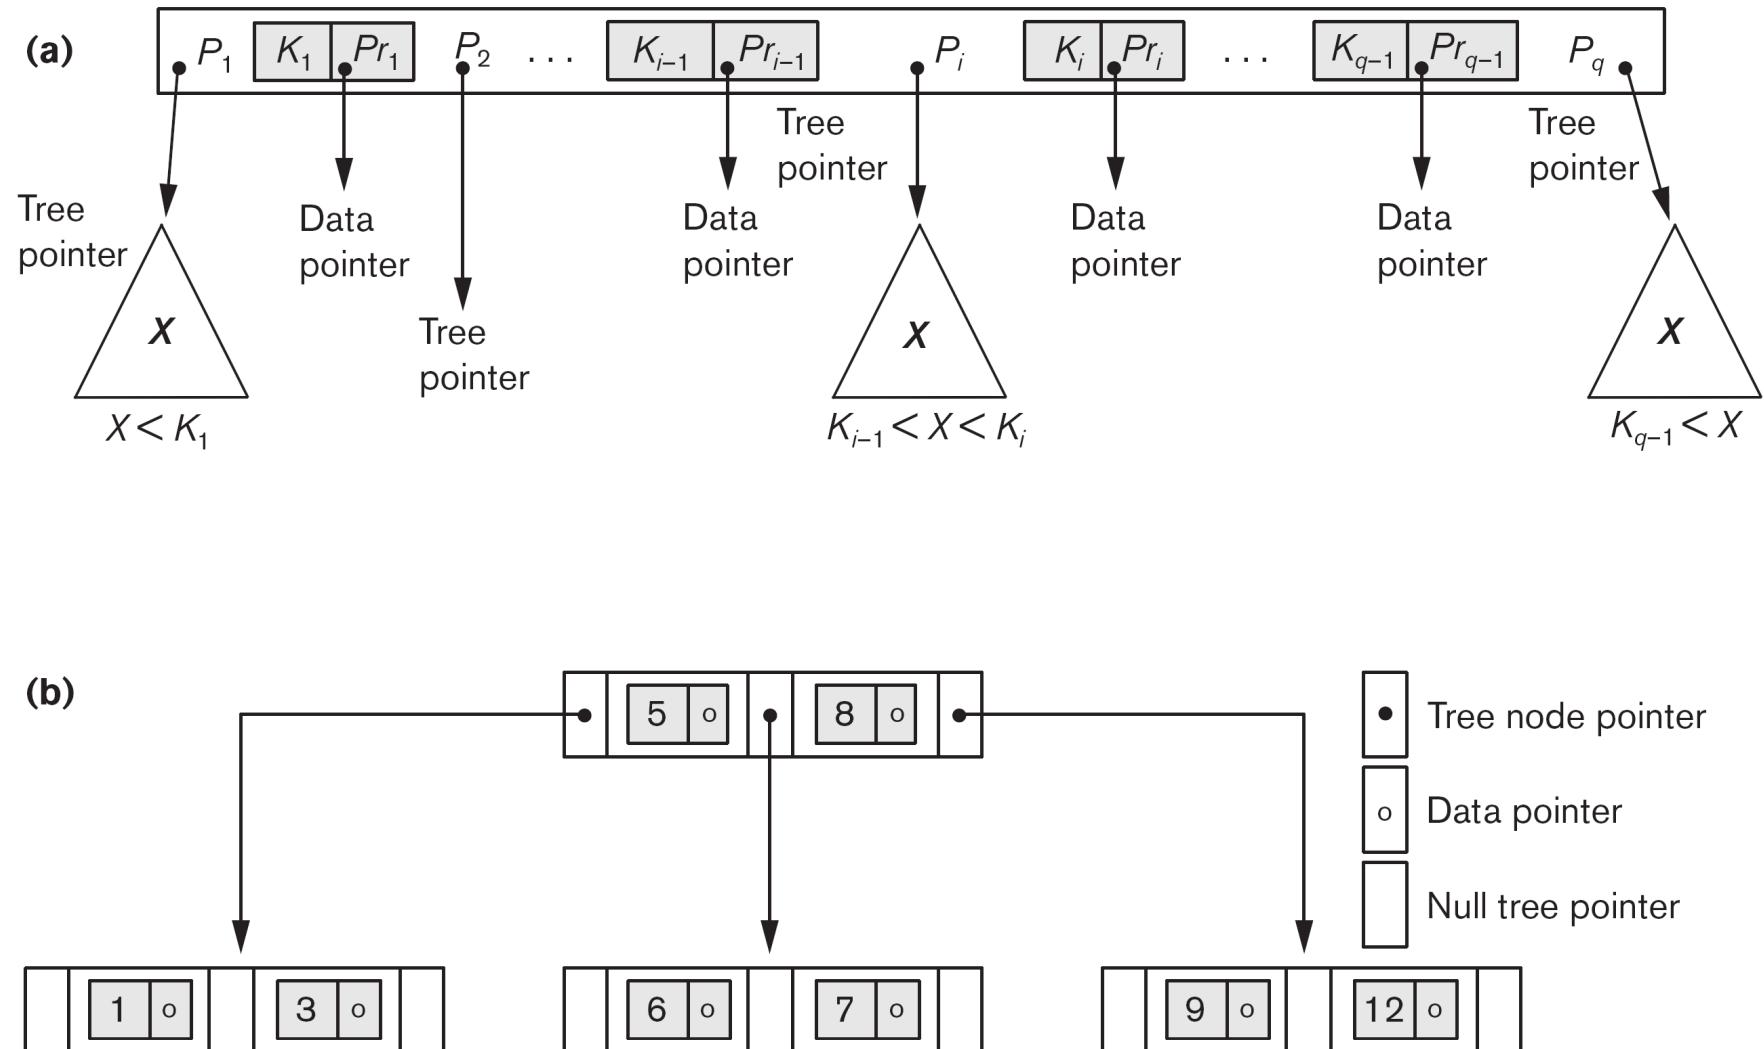


Figure 18.10

B-tree structures. (a) A node in a B-tree with $q - 1$ search values. (b) A B-tree of order $p = 3$. The values were inserted in the order 8, 5, 1, 7, 3, 12, 9, 6.

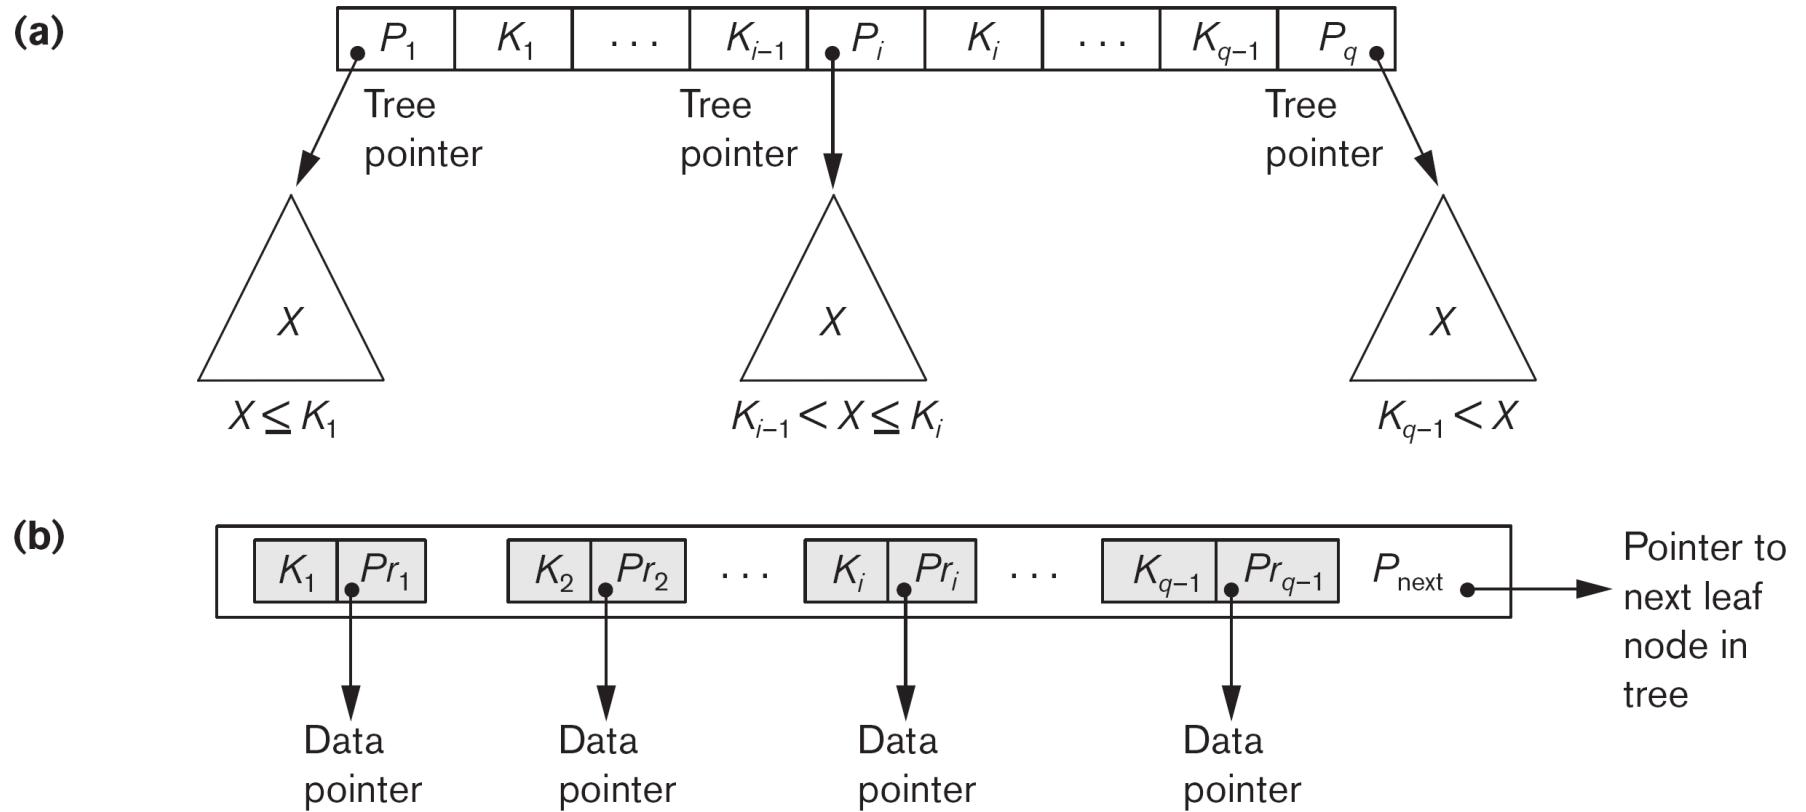
B⁺-Tree Example

The Nodes of a B⁺-tree

Figure 18.11

The nodes of a B⁺-tree. (a) Internal node of a B⁺-tree with $q - 1$ search values.

(b) Leaf node of a B⁺-tree with $q - 1$ search values and $q - 1$ data pointers.



Example of an Insertion in a B+-tree

- An example of insertion in a B+-tree with
 - $p = 3$
 - $p_{leaf} = 2$

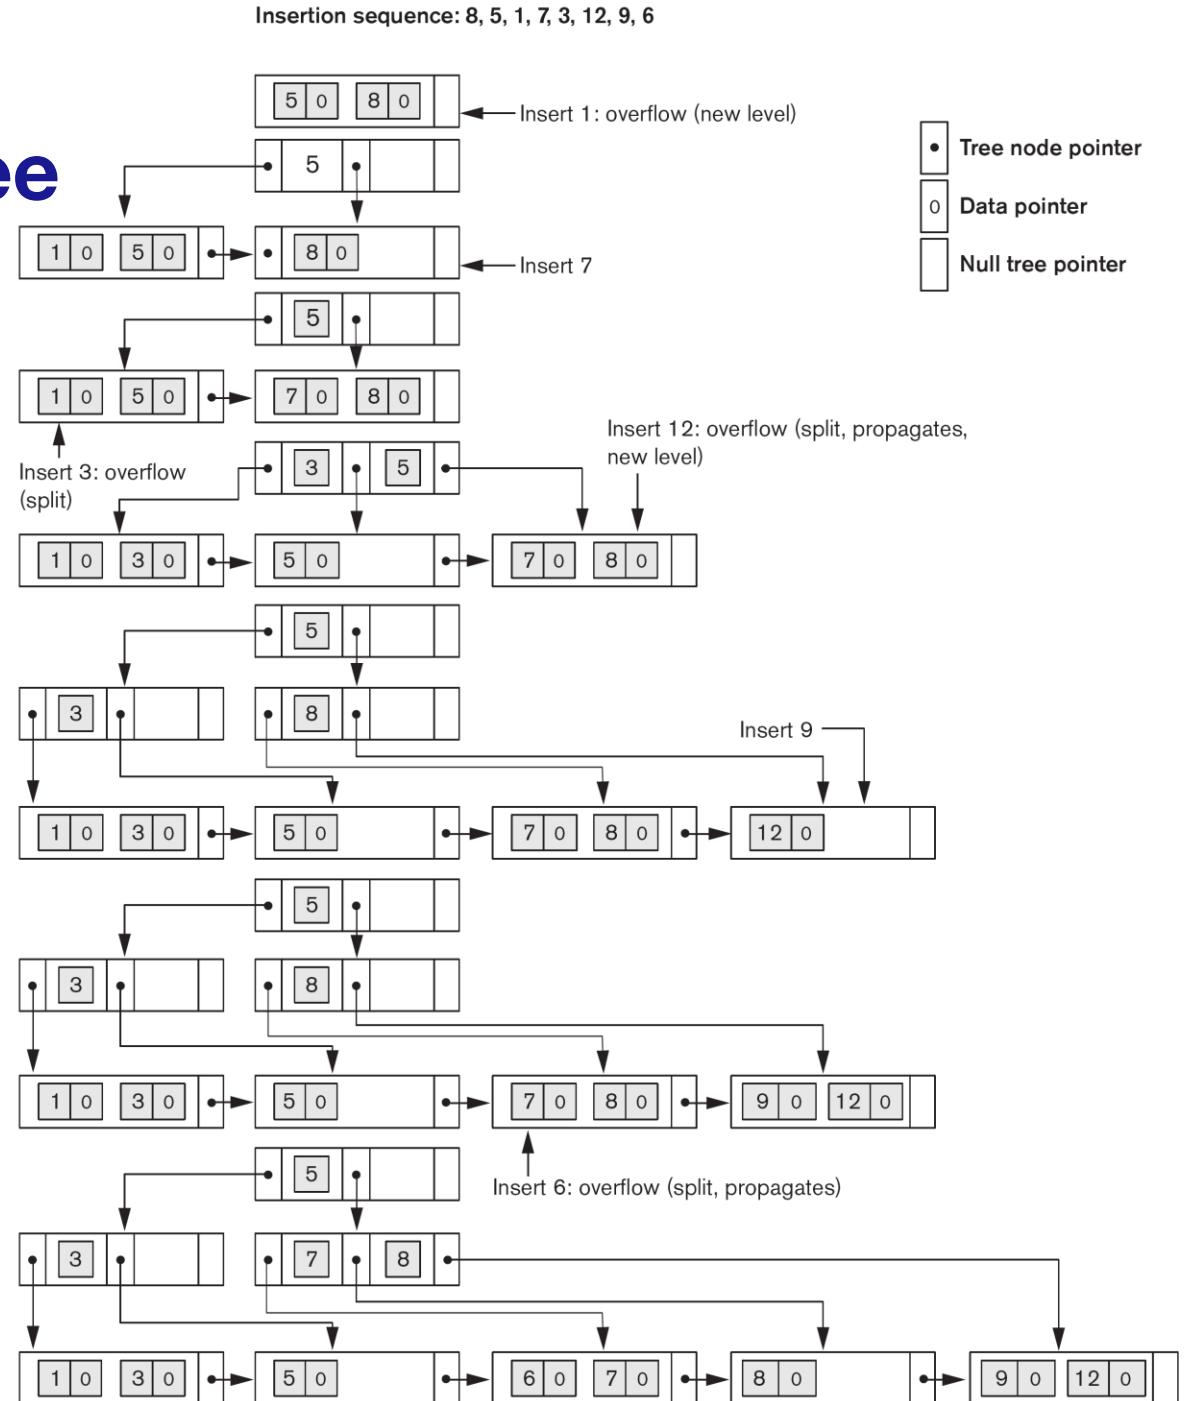
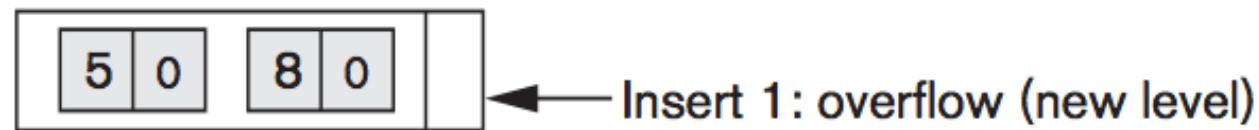


Figure 18.12

An example of insertion in a B+-tree with $p = 3$ and $p_{leaf} = 2$.

Example of an Insertion in a B⁺-tree

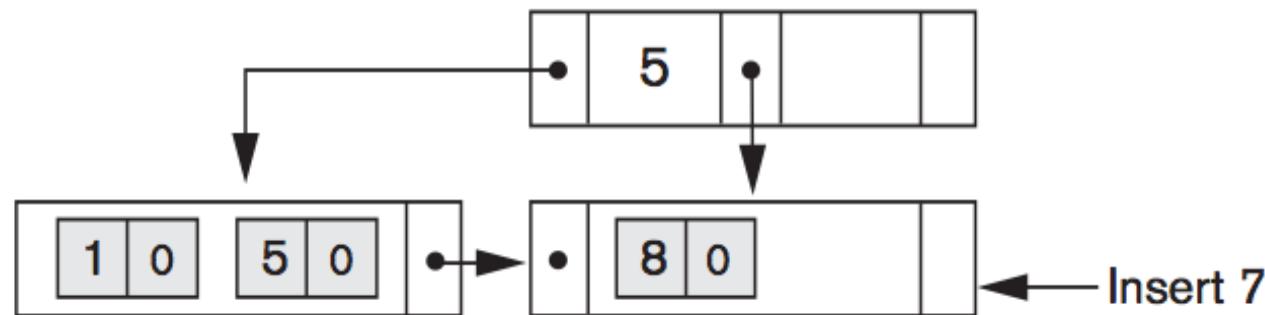
Insertion sequence: 8, 5, 1, 7, 3, 12, 9, 6



- Tree node pointer
- 0 Data pointer
- Null tree pointer

Example of an Insertion in a B⁺-tree

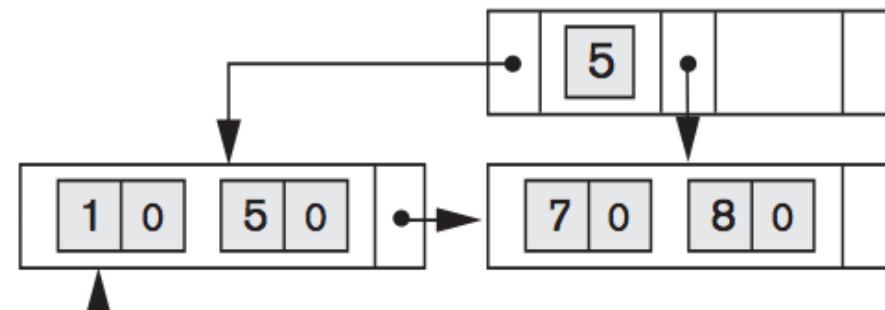
Insertion sequence: 8, 5, 1, 7, 3, 12, 9, 6



- Tree node pointer
- 0 Data pointer
- Null tree pointer

Example of an Insertion in a B⁺-tree

Insertion sequence: 8, 5, 1, 7, 3, 12, 9, 6

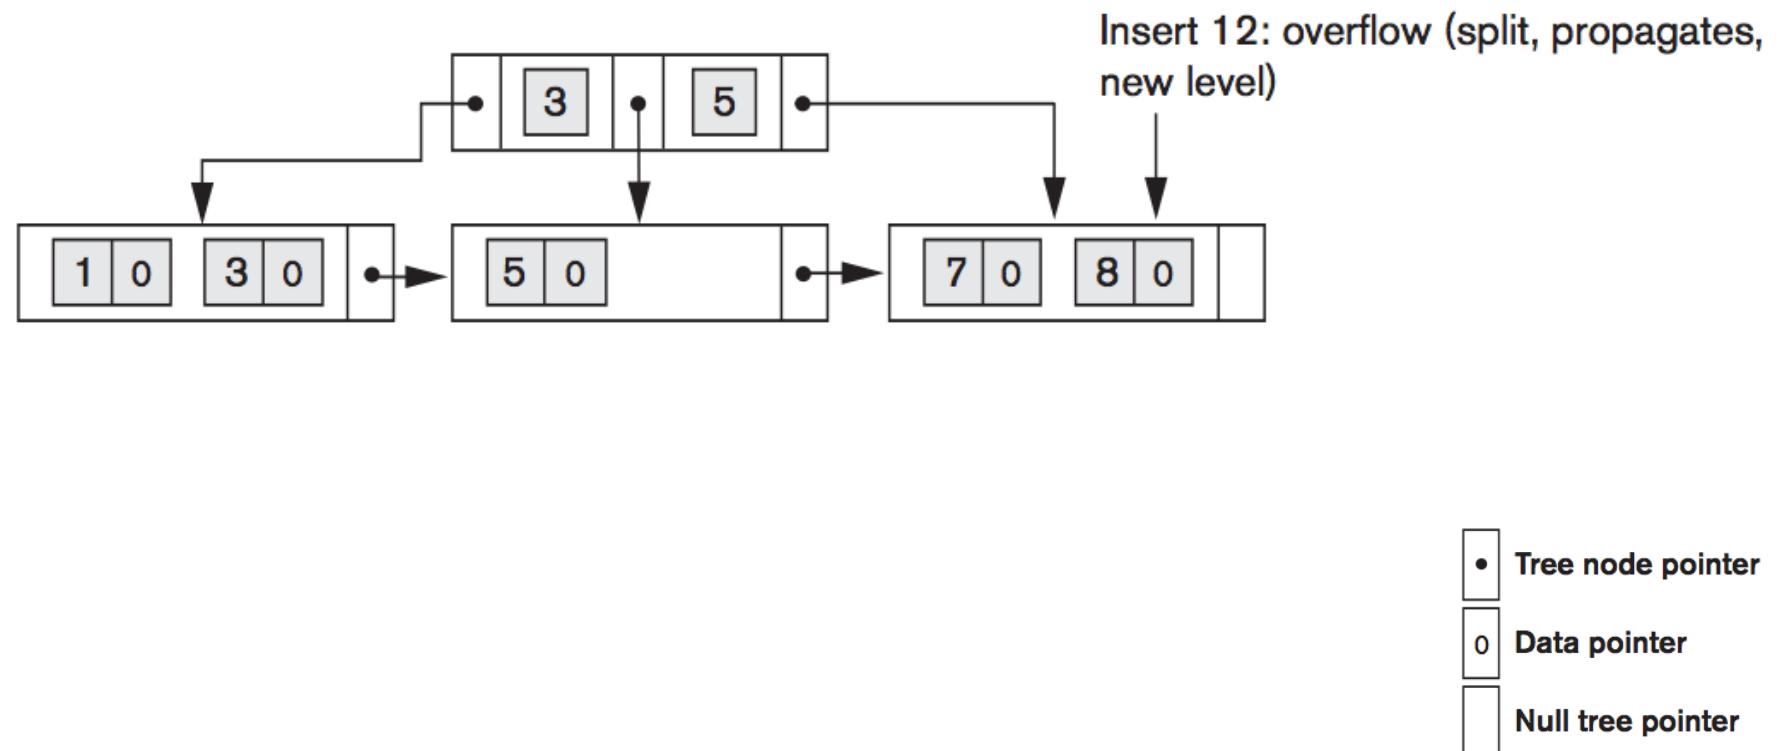


Insert 3: overflow
(split)

- Tree node pointer
- 0 Data pointer
- Null tree pointer

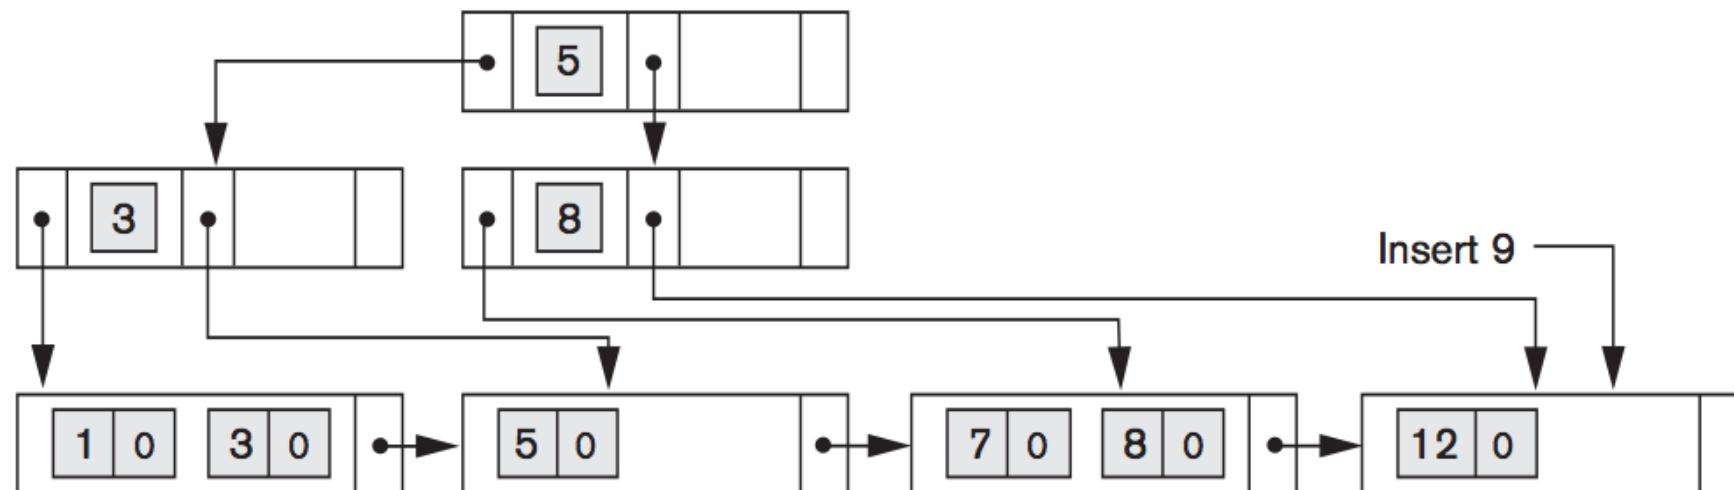
Example of an Insertion in a B⁺-tree

Insertion sequence: 8, 5, 1, 7, 3, 12, 9, 6



Example of an Insertion in a B⁺-tree

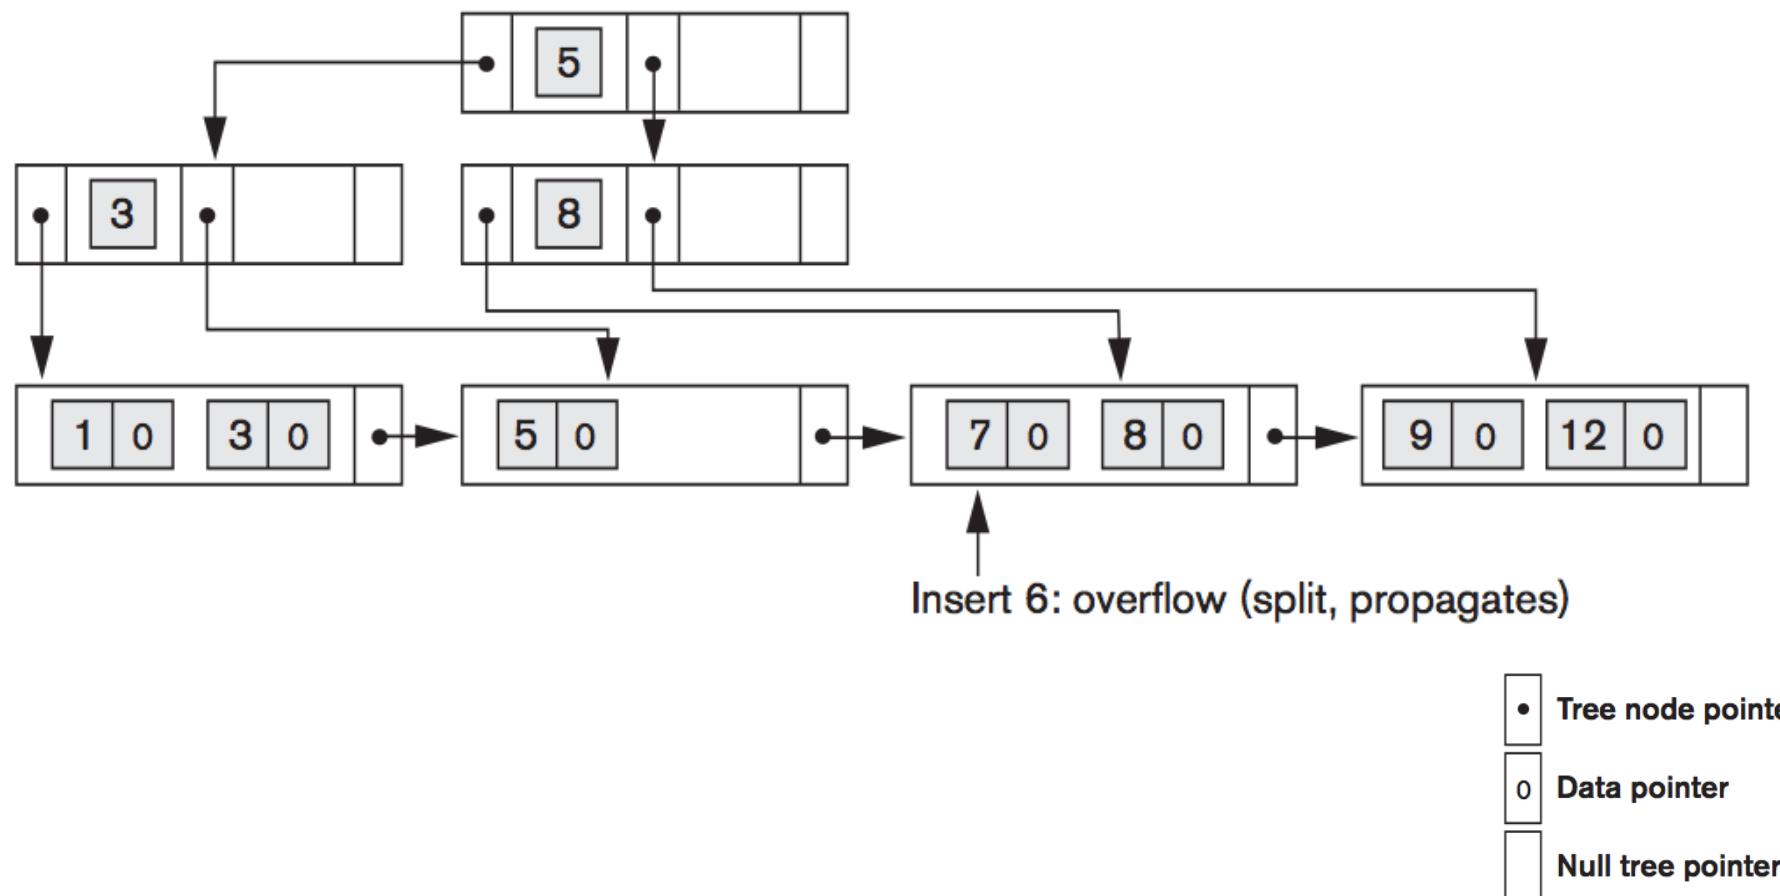
Insertion sequence: 8, 5, 1, 7, 3, 12, 9, 6



- Tree node pointer
- 0 Data pointer
- Null tree pointer

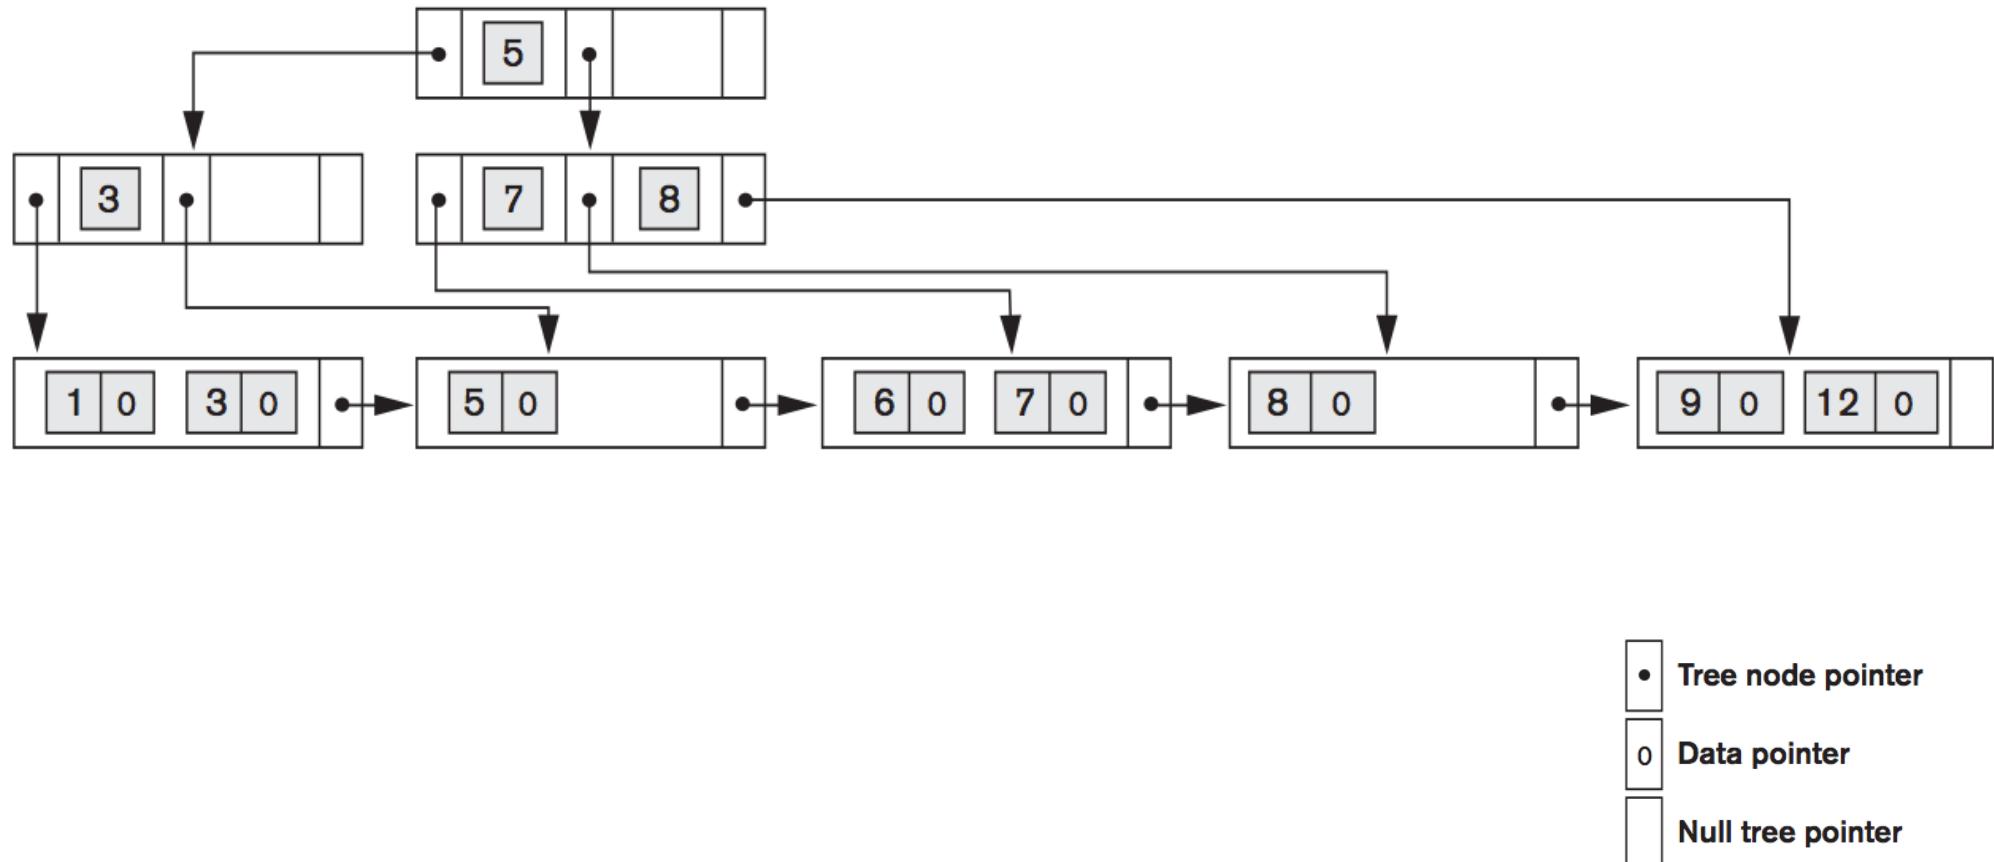
Example of an Insertion in a B⁺-tree

Insertion sequence: 8, 5, 1, 7, 3, 12, 9, 6



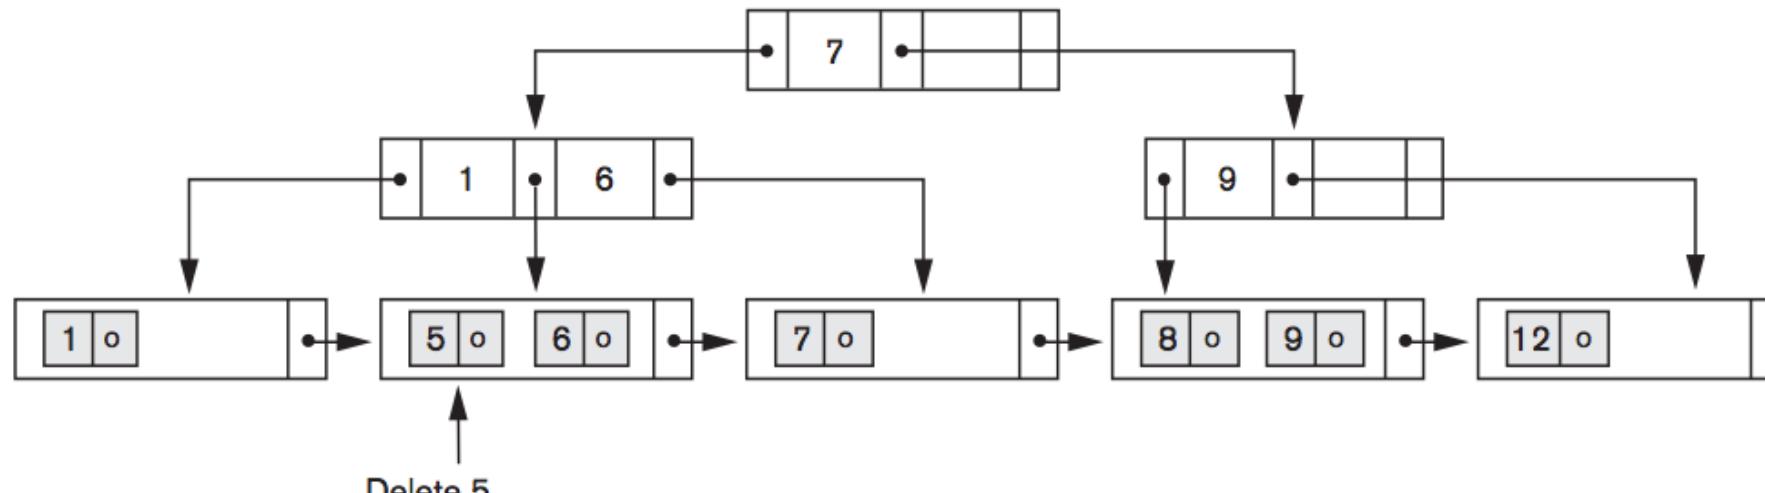
Example of an Insertion in a B⁺-tree

Insertion sequence: 8, 5, 1, 7, 3, 12, 9, 6



Example of a Deletion in a B⁺-tree

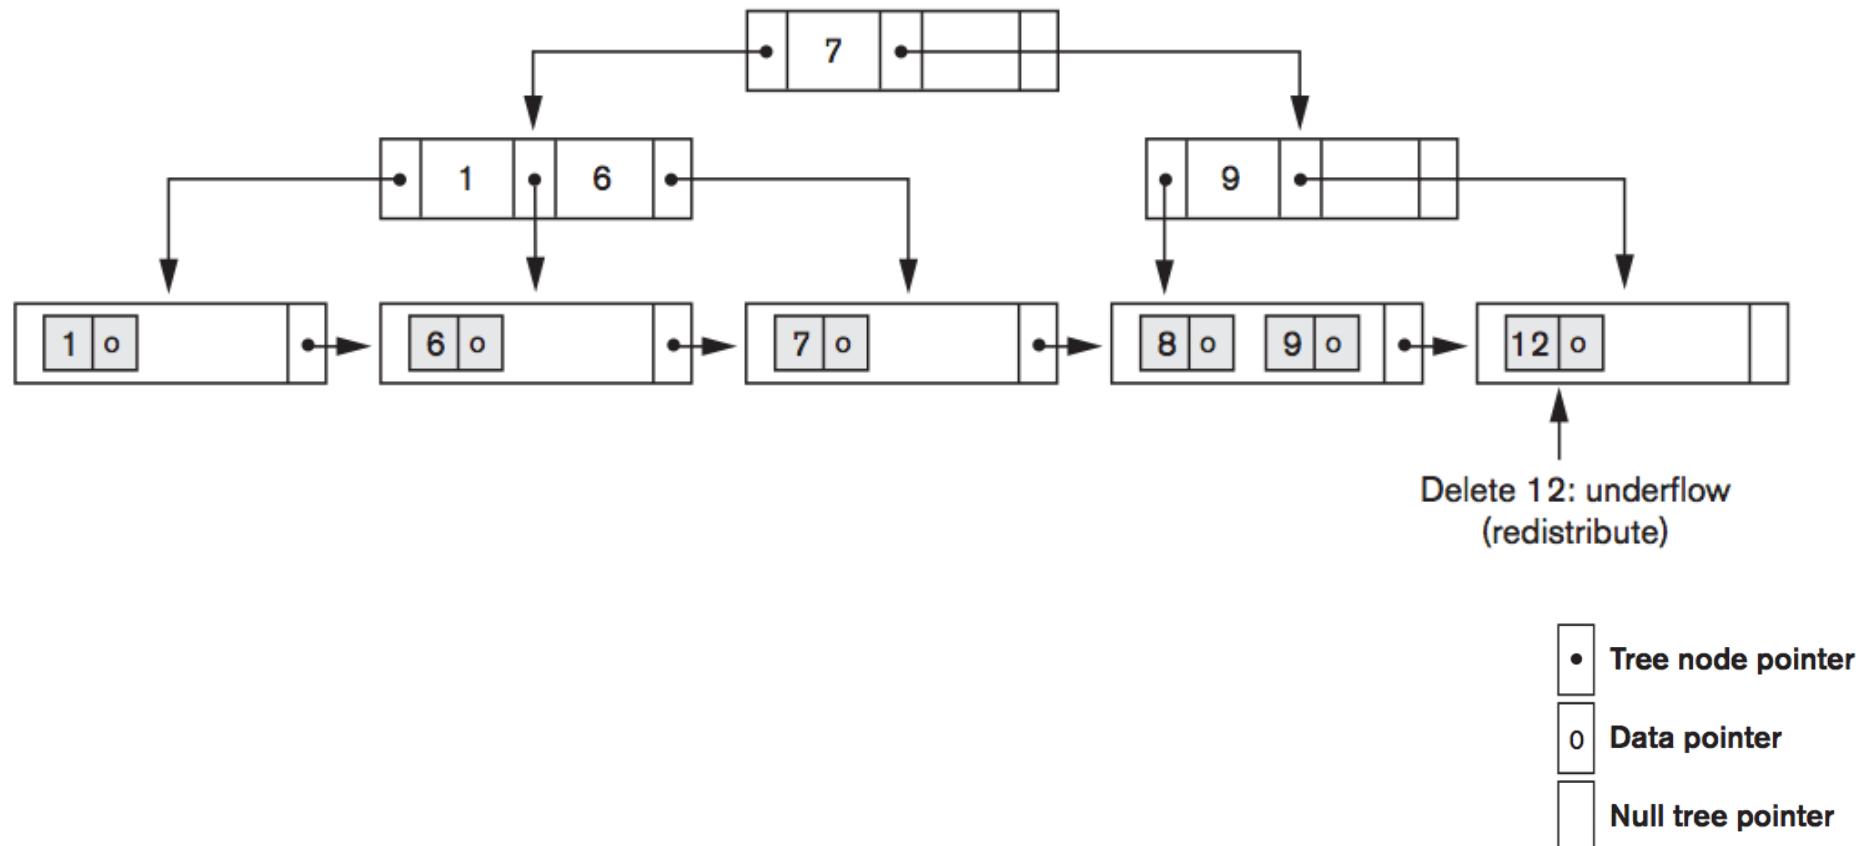
Deletion sequence: 5, 12, 9



- Tree node pointer
- 0 Data pointer
- Null tree pointer

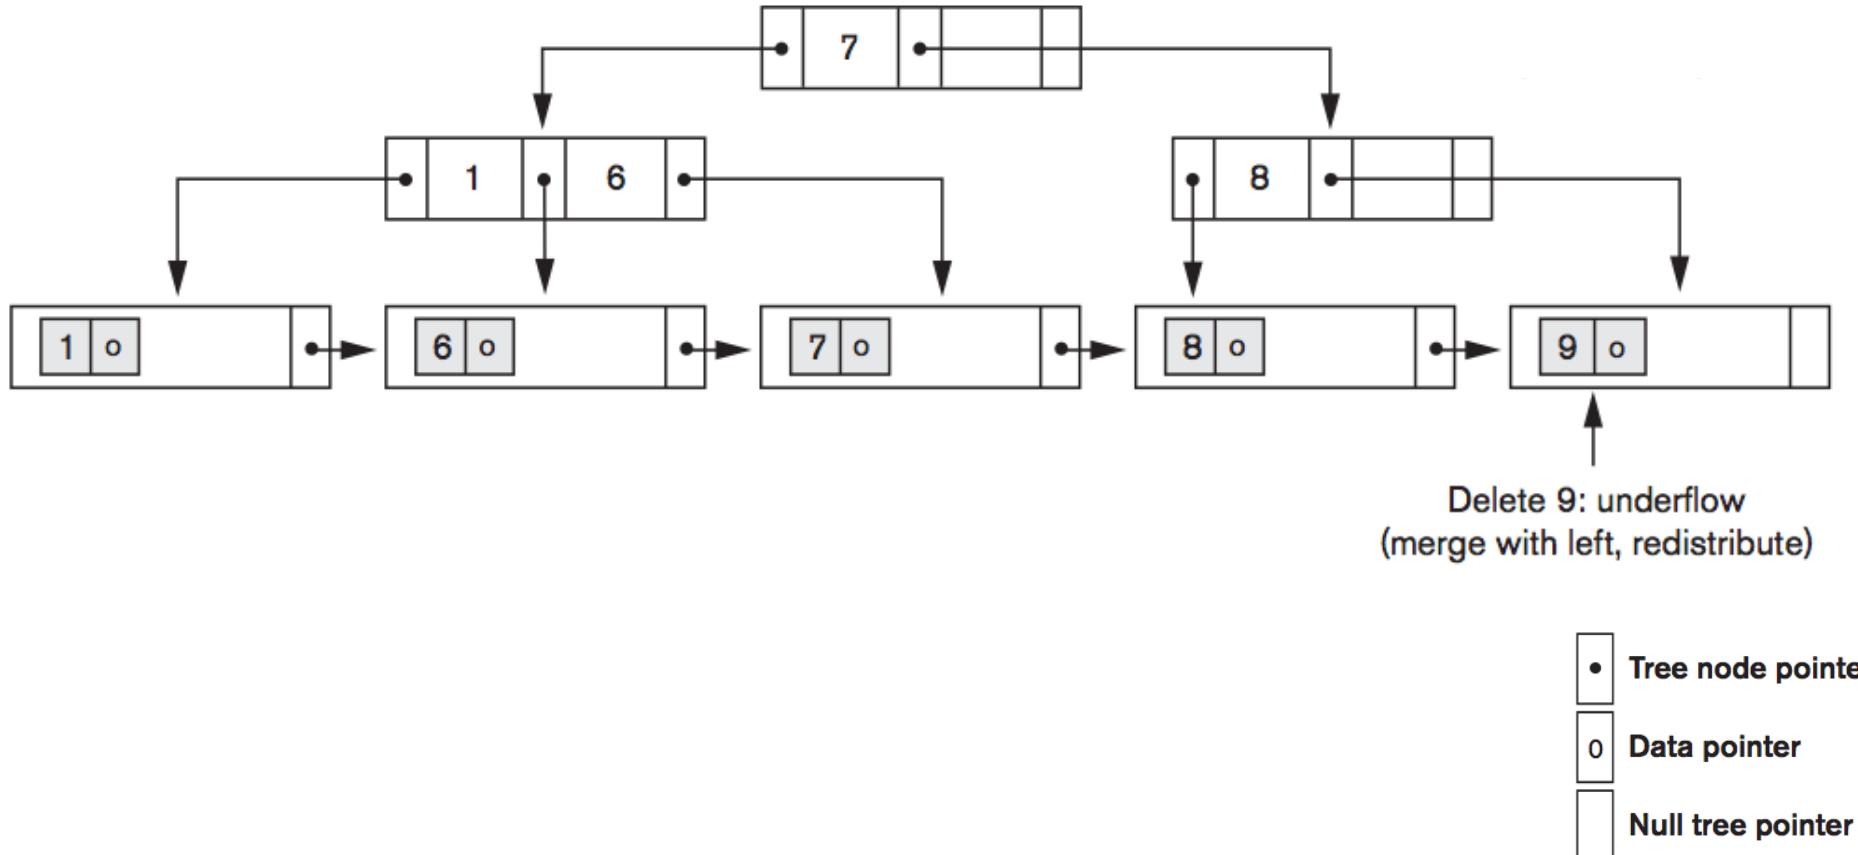
Example of a Deletion in a B⁺-tree

Deletion sequence: 5, 12, 9



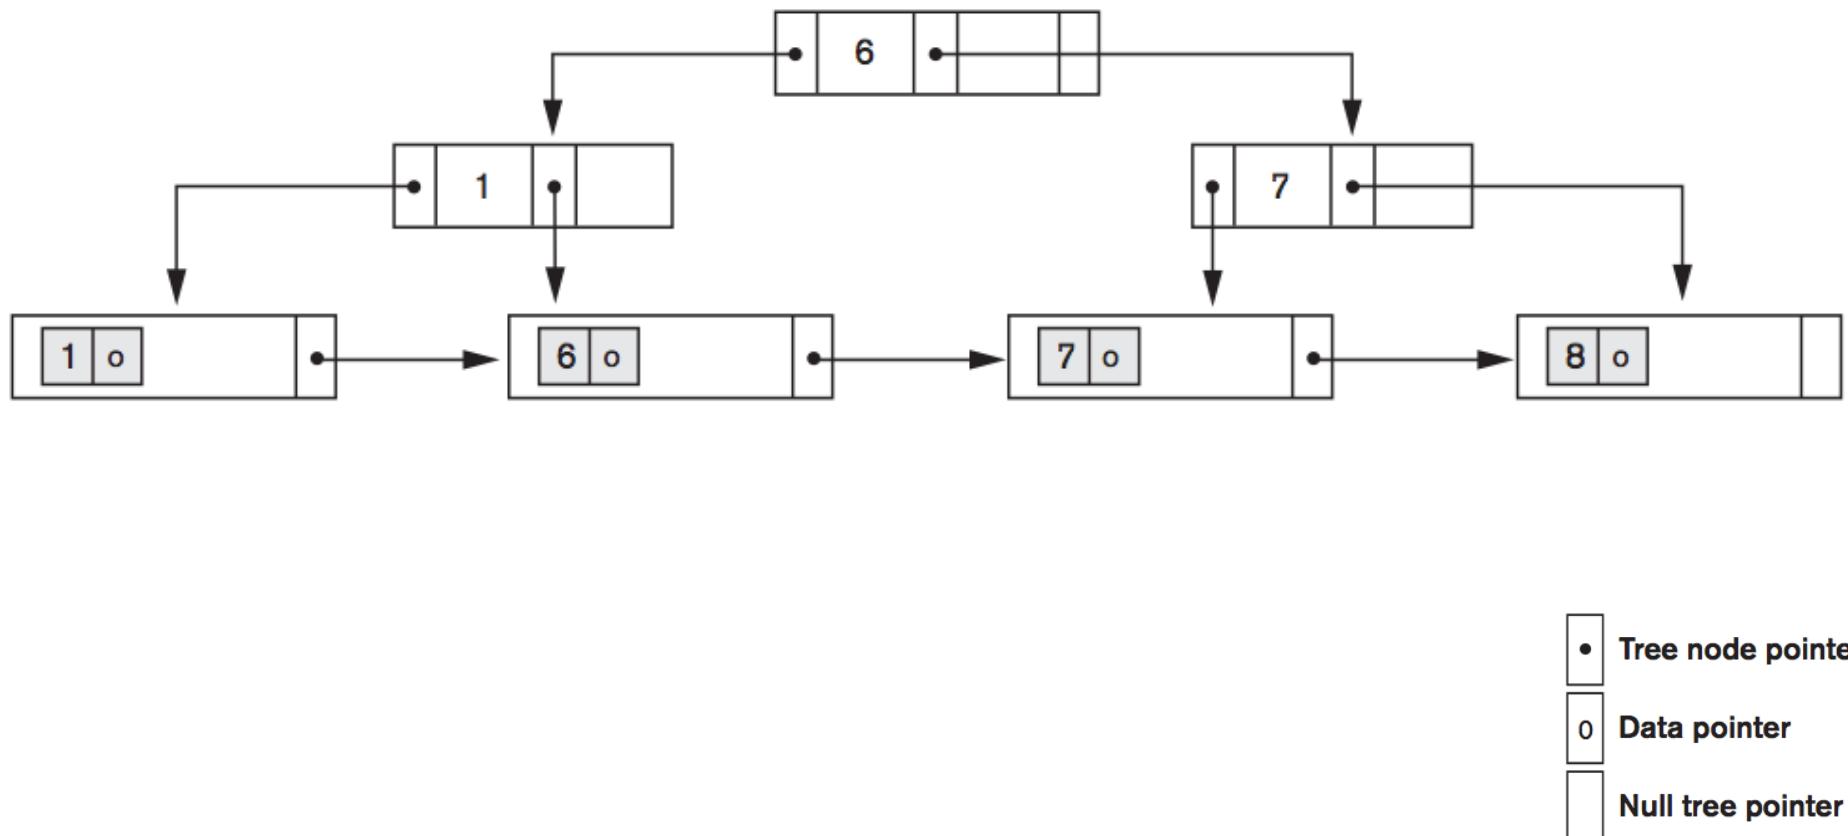
Example of a Deletion in a B⁺-tree

Deletion sequence: 5, 12, 9



Example of a Deletion in a B⁺-tree

Deletion sequence: 5, 12, 9



Summary

- Types of Single-level Ordered Indexes
 - Primary Indexes
 - Clustering Indexes
 - Secondary Indexes
- Multilevel Indexes
- Dynamic Multilevel Indexes Using B-Trees and B+-Trees